
MODULE *Config*

INSTANCE *Naturals*

INSTANCE *FiniteSets*

INSTANCE *Sequences*

INSTANCE *TLC*

An empty constant

CONSTANT *Nil*

Transaction status constants

CONSTANTS

TransactionPending,
TransactionValidating,
TransactionApplying,
TransactionComplete,
TransactionFailed

Configuration status constants

CONSTANTS

ConfigurationPending,
ConfigurationInitializing,
ConfigurationUpdating,
ConfigurationComplete,
ConfigurationFailed

The set of all nodes

CONSTANT *Node*

Target is the possible targets, paths, and values

Example: $Target \triangleq$ [
 $target1 \mapsto$ [
 $path1 \mapsto \{“value1”, “value2”\}$,
 $path2 \mapsto \{“value2”, “value3”\}$],
 $target2 \mapsto$ [
 $path2 \mapsto \{“value3”, “value4”\}$,
 $path3 \mapsto \{“value4”, “value5”\}$]]

CONSTANT *Target*

ASSUME *Nil* ∈ STRING

ASSUME *TransactionPending* ∈ STRING

ASSUME *TransactionValidating* ∈ STRING

ASSUME *TransactionApplying* ∈ STRING

ASSUME $TransactionComplete \in \text{STRING}$
 ASSUME $TransactionFailed \in \text{STRING}$

 ASSUME $ConfigurationPending \in \text{STRING}$
 ASSUME $ConfigurationInitializing \in \text{STRING}$
 ASSUME $ConfigurationUpdating \in \text{STRING}$
 ASSUME $ConfigurationComplete \in \text{STRING}$
 ASSUME $ConfigurationFailed \in \text{STRING}$

 ASSUME $\wedge IsFiniteSet(Node)$
 $\wedge \forall n \in Node :$
 $\wedge n \notin \text{DOMAIN } Target$
 $\wedge n \in \text{STRING}$

 ASSUME $\wedge \forall t \in \text{DOMAIN } Target :$
 $\wedge t \notin Node$
 $\wedge t \in \text{STRING}$
 $\wedge \forall p \in \text{DOMAIN } Target[t] :$
 $IsFiniteSet(Target[t][p])$

```

TYPE TransactionStatus ::= status ∈
  { TransactionPending,
    TransactionValidating,
    TransactionApplying,
    TransactionComplete,
    TransactionFailed }

TYPE Transaction ≜ [
  id      ::= id ∈ STRING,
  index   ::= index ∈ Nat,
  revision ::= revision ∈ Nat,
  atomic  ::= atomic ∈ BOOLEAN ,
  sync    ::= sync ∈ BOOLEAN ,
  changes ::= [ target ∈ SUBSET (DOMAIN Target) ↦ [
    path ∈ SUBSET (DOMAIN Target[target]) ↦ [
      value ::= value ∈ STRING,
      delete ::= delete ∈ BOOLEAN ]],
  status ::= status ∈ TransactionStatus ]

TYPE ConfigurationStatus ::= status ∈
  { ConfigurationPending,
    ConfigurationInitializing,
    ConfigurationUpdating,
    ConfigurationComplete,
    ConfigurationFailed }

TYPE Configuration ≜ [
  id      ::= id ∈ STRING,
  revision ::= revision ∈ Nat,

```

```

target ::= target ∈ STRING,
paths ::= [ path ∈ SUBSET (DOMAIN Target[target]) ↦ [
    value ::= value ∈ STRING,
    index ::= index ∈ Nat,
    deleted ::= delete ∈ BOOLEAN ]],
txIndex ::= txIndex ∈ Nat,
syncIndex ::= syncIndex ∈ Nat,
term      ::= term ∈ Nat,
status ::= status ∈ ConfigurationStatus]

```

A sequence of transactions

Each transactions contains a record of 'changes' for a set of targets

VARIABLE *transaction*

A record of target configurations

Each configuration represents the desired state of the target

VARIABLE *configuration*

A record of target states

VARIABLE *target*

A record of target masters

VARIABLE *master*

VARIABLE *history*

$\text{vars} \triangleq \langle \text{transaction}, \text{configuration}, \text{master}, \text{target}, \text{history} \rangle$

$\text{ChangeMaster}(n, t) \triangleq$
 $\wedge \text{master}[t].\text{master} \neq n$
 $\wedge \text{master}' = [\text{master} \text{ EXCEPT } ![t].\text{term} = \text{master}[t].\text{term} + 1,$
 $\phantom{\wedge \text{master}' = } ![t].\text{master} = n]$
 $\wedge \text{UNCHANGED } \langle \text{transaction}, \text{configuration}, \text{target}, \text{history} \rangle$

This section models the northbound *API* for the configuration service.

$\text{ValidValues}(t, p) \triangleq$
 $\text{UNION } \{ \{ [value \mapsto v, delete \mapsto \text{FALSE}] : v \in \text{Target}[t][p] \}, \{ [value \mapsto \text{Nil}, delete \mapsto \text{TRUE}] \} \}$

$\text{ValidPaths}(t) \triangleq$
 $\text{UNION } \{ \{ v @@@ [path \mapsto p] : v \in \text{ValidValues}(t, p) \} : p \in \text{DOMAIN Target}[t] \}$

$\text{ValidTargets} \triangleq$
 $\text{UNION } \{ \{ p @@@ [target \mapsto t] : p \in \text{ValidPaths}(t) \} : t \in \text{DOMAIN Target} \}$

$\text{ValidPath}(s, t, p) \triangleq$
 $\text{LET } value \triangleq \text{CHOOSE } v \in s : v.\text{target} = t \wedge v.\text{path} = p$

IN
 $[value \mapsto value.value,$
 $delete \mapsto value.delete]$

 $ValidTarget(s, t) \triangleq$
 $[p \in \{v.path : v \in \{v \in s : v.target = t\}\} \mapsto ValidPath(s, t, p)]$

 $ValidChange(s) \triangleq$
 $[t \in \{v.target : v \in s\} \mapsto ValidTarget(s, t)]$

 $ValidChanges \triangleq$
 LET $changeSets \triangleq \{s \in \text{SUBSET } ValidTargets :$
 $\forall t \in \text{DOMAIN } Target :$
 $\forall p \in \text{DOMAIN } Target[t] :$
 $Cardinality(\{v \in s : v.target = t \wedge v.path = p\}) \leq 1\}$

 IN
 $\{ValidChange(s) : s \in changeSets\}$

 $NextIndex \triangleq$
 IF $Len(transaction) = 0$ THEN
 1
 ELSE
 LET $i \triangleq \text{CHOOSE } i \in \text{DOMAIN } transaction :$
 $\forall j \in \text{DOMAIN } transaction :$
 $transaction[j].index \leq transaction[i].index$
 IN $transaction[i].index + 1$

 Add a set of changes to the transaction log
 $Change \triangleq$
 $\wedge \exists changes \in ValidChanges :$
 $\wedge transaction' = Append(transaction, [index \mapsto NextIndex,$
 $atomic \mapsto \text{FALSE},$
 $sync \mapsto \text{FALSE},$
 $changes \mapsto changes,$
 $status \mapsto TransactionPending])$
 $\wedge \text{UNCHANGED } \langle configuration, master, target, history \rangle$

This section models the Transaction log reconciler.

Reconcile the transaction log
 $ReconcileTransaction(n, t) \triangleq$
 If the transaction is Pending, begin validation if the prior transaction
 has already been applied. This simplifies concurrency control in the controller
 and guarantees transactions are applied to the configurations in sequential order.
 $\wedge \forall \wedge transaction[t].status = TransactionPending$

$\wedge \vee \wedge transaction[t].index - 1 \in \text{DOMAIN } transaction$
 $\wedge transaction[transaction[t].index - 1].status \in \{TransactionComplete, TransactionFailed\}$
 $\vee transaction[t].index - 1 \notin \text{DOMAIN } transaction$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionValidating]$
 $\wedge \text{UNCHANGED } \langle configuration, history \rangle$
 If the transaction is in the Validating state, compute and validate the
 Configuration for each target.
 $\vee \wedge transaction[t].status = TransactionValidating$
 If validation fails any target, mark the transaction Failed.
 If validation is successful, proceed to Applying.
 $\wedge \exists valid \in \text{BOOLEAN} :$
 $\vee \wedge valid$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionApplying]$
 $\vee \wedge \neg valid$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionFailed]$
 $\wedge \text{UNCHANGED } \langle configuration, history \rangle$
 If the transaction is in the Applying state, update the Configuration for each
 target and Complete the transaction.
 $\vee \wedge transaction[t].status = TransactionApplying$
 $\wedge \vee \wedge transaction[t].atomic$
 TODO: Apply atomic transactions here
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionComplete]$
 $\wedge \text{UNCHANGED } \langle configuration, history \rangle$
 $\vee \wedge \neg transaction[t].atomic$
 Add the transaction index to each updated path
 $\wedge configuration' = [$
 $\quad r \in \text{DOMAIN } Target \mapsto$
 $\quad \text{IF } r \in \text{DOMAIN } transaction[t].changes \text{ THEN}$
 $\quad \quad [configuration[r] \text{ EXCEPT}$
 $\quad \quad \quad !.paths = [p \in \text{DOMAIN } transaction[t].changes[r] \mapsto$
 $\quad \quad \quad \quad [index \mapsto transaction[t].index,$
 $\quad \quad \quad \quad \quad value \mapsto transaction[t].changes[r][p].value,$
 $\quad \quad \quad \quad \quad deleted \mapsto transaction[t].changes[r][p].delete]]$
 $\quad \quad \quad @@ configuration[r].paths,$
 $\quad \quad \quad !.txIndex = transaction[t].index,$
 $\quad \quad \quad !.status = ConfigurationPending]$
 $\quad \text{ELSE}$
 $\quad \quad configuration[r]]$
 $\wedge history' = [r \in \text{DOMAIN } Target \mapsto Append(history[r], configuration'[r])]$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionComplete]$
 $\wedge \text{UNCHANGED } \langle master, target \rangle$

This section models the Configuration reconciler.

$ReconcileConfiguration(n, c) \triangleq$
 $\wedge \vee \wedge configuration[c].status = ConfigurationPending$
 $\wedge master[configuration[c].target].master \neq Nil$
 If the configuration is marked *ConfigurationPending* and mastership has changed (indicated by an increased mastership term), mark the configuration *ConfigurationInitializing* to force full re-synchronization.
 $\wedge \vee \wedge master[configuration[c].target].term > configuration[c].term$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationInitializing,$
 $![c].term = master[configuration[c].target].term]$
 $\wedge history' = [history \text{ EXCEPT } ![c] = Append(history[c], configuration'[c])]$
 If the configuration is marked *ConfigurationPending* and the values have changed (determined by comparing the transaction index to the last *sync* index), mark the configuration *ConfigurationUpdating* to push the changes to the target.
 $\vee \wedge master[configuration[c].target].term = configuration[c].term$
 $\wedge configuration[c].syncIndex < configuration[c].txIndex$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationUpdating]$
 $\wedge history' = [history \text{ EXCEPT } ![c] = Append(history[c], configuration'[c])]$
 $\wedge UNCHANGED \langle target \rangle$
 $\vee \wedge configuration[c].status = ConfigurationInitializing$
 $\wedge master[configuration[c].target].master = n$
 Merge the configuration paths with the target paths, removing paths that have been marked deleted
 $\wedge LET \ deletePaths \triangleq \{p \in DOMAIN \ configuration[c].paths : configuration[c].paths[p].deleted\}$
 $\quad \quad \quad configPaths \triangleq DOMAIN \ configuration[c].paths \setminus deletePaths$
 $\quad \quad \quad targetPaths \triangleq DOMAIN \ target[configuration[c].target] \setminus deletePaths$
 IN
 $\wedge target' = [target \text{ EXCEPT } ![configuration[c].target] =$
 $\quad [p \in configPaths \mapsto [value \mapsto configuration[c].paths[p]]]$
 $\quad @@ [p \in targetPaths \mapsto target[configuration[c].target][p]]]$
 Set the configuration's status to Complete
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationComplete,$
 $\quad \quad \quad ![c].syncIndex = configuration[c].txIndex]$
 $\wedge history' = [history \text{ EXCEPT } ![c] = Append(history[c], configuration'[c])]$
 If the configuration is marked *ConfigurationUpdating*, we only need to push paths that have changed since the target was initialized or last updated by the controller. The set of changes made since the last synchronization are identified by comparing the index of each path-value to the last synchronization index, *syncIndex*
 $\vee \wedge configuration[c].status = ConfigurationUpdating$
 $\wedge master[configuration[c].target].master = n$
 Compute the set of updated and deleted paths by comparing their indexes to the target's last sync index.
 $\wedge LET \ updatePaths \triangleq \{p \in DOMAIN \ configuration[c].paths :$
 $\quad \quad \quad configuration[c].paths[p].index > configuration[c].syncIndex\}$

$$\begin{aligned}
deletePaths &\triangleq \{p \in updatePaths : configuration[c].paths[p].deleted\} \\
configPaths &\triangleq updatePaths \setminus deletePaths \\
targetPaths &\triangleq \text{DOMAIN } target[configuration[c].target] \setminus deletePaths
\end{aligned}$$

IN

Update the target paths by adding/updating paths that have changed and removing paths that have been deleted since the last *sync*.

$$\begin{aligned}
&\wedge target' = [target \text{ EXCEPT } ![configuration[c].target] = \\
&\quad [p \in configPaths \mapsto configuration[c].paths[p]] \\
&\quad @@ [p \in targetPaths \mapsto target[configuration[c].target][p]] \\
&\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationComplete, \\
&\quad ![c].syncIndex = configuration[c].txIndex] \\
&\wedge history' = [history \text{ EXCEPT } ![c] = Append(history[c], configuration'[c])]
\end{aligned}$$

If the configuration is not already *ConfigurationPending* and mastership has been lost revert it. This can occur when the connection to the target has been lost and the mastership is no longer valid.

TODO: We still need to model mastership changes

$$\begin{aligned}
&\vee \wedge configuration[c].status \neq ConfigurationPending \\
&\quad \wedge master[configuration[c].target].master = Nil \\
&\quad \wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationPending] \\
&\quad \wedge history' = [history \text{ EXCEPT } ![c] = Append(history[c], configuration'[c])] \\
&\quad \wedge \text{UNCHANGED } \langle target \rangle \\
&\wedge \text{UNCHANGED } \langle transaction, master \rangle
\end{aligned}$$

Init and next state predicates

$$\begin{aligned}
Init &\triangleq \\
&\wedge transaction = \langle \rangle \\
&\wedge configuration = [t \in \text{DOMAIN } Target \mapsto \\
&\quad [target \mapsto t, \\
&\quad paths \mapsto \\
&\quad \quad [path \in \{\} \mapsto \\
&\quad \quad \quad [path \mapsto path, \\
&\quad \quad \quad \quad value \mapsto Nil, \\
&\quad \quad \quad \quad index \mapsto 0, \\
&\quad \quad \quad \quad deleted \mapsto FALSE]], \\
&\quad txIndex \mapsto 0, \\
&\quad syncIndex \mapsto 0, \\
&\quad term \mapsto 0, \\
&\quad status \mapsto ConfigurationPending]] \\
&\wedge target = [t \in \text{DOMAIN } Target \mapsto \\
&\quad [path \in \{\} \mapsto \\
&\quad \quad [value \mapsto Nil]]] \\
&\wedge master = [t \in \text{DOMAIN } Target \mapsto [master \mapsto Nil, term \mapsto 0]] \\
&\wedge history = [t \in \text{DOMAIN } Target \mapsto \langle \rangle]
\end{aligned}$$

$$\begin{aligned}
Next &\triangleq \\
&\vee \textit{Change} \\
&\vee \exists n \in \textit{Node} : \\
&\quad \vee \exists t \in \textit{DOMAIN Target} : \\
&\quad \quad \textit{ChangeMaster}(n, t) \\
&\quad \vee \exists t \in \textit{DOMAIN transaction} : \\
&\quad \quad \textit{ReconcileTransaction}(n, t) \\
&\quad \vee \exists t \in \textit{DOMAIN configuration} : \\
&\quad \quad \textit{ReconcileConfiguration}(n, t) \\
Spec &\triangleq \textit{Init} \wedge \Box [Next]_{vars} \\
Inv &\triangleq \\
&\wedge \forall a, b \in \textit{DOMAIN transaction} : \\
&\quad \textit{transaction}[a].\textit{index} > \textit{transaction}[b].\textit{index} \Rightarrow \\
&\quad (\textit{transaction}[a].\textit{status} \in \{\textit{TransactionComplete}, \textit{TransactionFailed}\} \Rightarrow \\
&\quad \quad \textit{transaction}[b].\textit{status} \in \{\textit{TransactionComplete}, \textit{TransactionFailed}\}) \\
&\wedge \forall t \in \textit{DOMAIN Target} : \\
&\quad \forall c \in \textit{DOMAIN history}[t] : \\
&\quad \quad \wedge \textit{configuration}[t].\textit{txIndex} \geq \textit{history}[t][c].\textit{txIndex} \\
&\quad \quad \wedge \textit{configuration}[t].\textit{syncIndex} \geq \textit{history}[t][c].\textit{syncIndex}
\end{aligned}$$

\ * Modification History
\ * Last modified Tue Jan 18 12:53:43 PST 2022 by jordanhalterman
\ * Created Wed Sep 22 13:22:32 PDT 2021 by jordanhalterman