

---

MODULE *Config*

---

INSTANCE *Naturals*  
 INSTANCE *FiniteSets*  
 INSTANCE *Sequences*  
 INSTANCE *TLC*

---

An empty constant  
 CONSTANT *Nil*

Transaction type constants  
 CONSTANTS  
     *Change*,  
     *Rollback*

Transaction isolation constants  
 CONSTANTS  
     *ReadCommitted*,  
     *Serializable*

Phase constants  
 CONSTANTS  
     *Initialize*,  
     *Validate*,  
     *Abort*,  
     *Commit*,  
     *Apply*

$Phase \triangleq$   
 LET  $phases \triangleq \langle$  *Initialize*,  
                     *Validate*,  
                     *Abort*,  
                     *Commit*,  
                     *Apply* $\rangle$   
 IN    $[p \in \{phases[i] : i \in \text{DOMAIN } phases\} \mapsto$   
       CHOOSE  $i \in \text{DOMAIN } phases : phases[i] = p]$

Status constants  
 CONSTANTS  
     *Pending*,  
     *Complete*,  
     *Failed*

$Status \triangleq$

```

LET statuses  $\triangleq$   $\langle$ Pending,
                  Complete,
                  Failed $\rangle$ 
IN  [s  $\in$  {statuses[i] : i  $\in$  DOMAIN statuses}  $\mapsto$ 
      CHOOSE i  $\in$  DOMAIN statuses : statuses[i] = s]

```

CONSTANTS

*Valid*,  
*Invalid*

CONSTANTS

*Success*,  
*Failure*

The set of all nodes

CONSTANT *Node*

Target is the set of all targets and their possible paths and values.

Example: *Target*  $\triangleq$  [  
     *target1*  $\mapsto$  [*persistent*  $\mapsto$  FALSE, *values*  $\mapsto$  [  
         *path1*  $\mapsto$  {"*value1*", "*value2*"},  
         *path2*  $\mapsto$  {"*value2*", "*value3*" }]],  
     *target2*  $\mapsto$  [*persistent*  $\mapsto$  TRUE, *values*  $\mapsto$  [  
         *path2*  $\mapsto$  {"*value3*", "*value4*" },  
         *path3*  $\mapsto$  {"*value4*", "*value5*" }]]]

CONSTANT *Target*

---

Configuration update/rollback requests are tracked and processed through two data types. Transactions represent the lifecycle of a single configuration change request and are stored in an append-only log. Configurations represent the desired configuration of a *gNMI* target based on the aggregate of relevant changes in the Transaction log.

TYPE *Type* ::= *type*  $\in$   
     {*Change*,  
     *Rollback*}

TYPE *Phase* ::= *phase*  $\in$   
     {*Initialize*,  
     *Validate*,  
     *Abort*,  
     *Commit*,  
     *Apply*}

TYPE *Status* ::= *status*  $\in$   
     {*Pending*,  
     *Initializing*,  
     *Initialized*,  
     *Validating*,  
     *Validated*,

```

    Committing,
    Committed,
    Applying,
    Applied,
    Synchronizing,
    Synchronized,
    Persisted,
    Failed}

TYPE Transaction  $\triangleq$  [
  type      ::= type  $\in$  Type,
  index     ::= index  $\in$  Nat,
  isolation ::= isolation  $\in$  {IsolationDefault, IsolationSerializable}
  values ::= [
    target  $\in$  SUBSET (DOMAIN Target)  $\mapsto$  [ path  $\in$  SUBSET (DOMAIN Target[target].values)  $\mapsto$ 
    [
      value ::= value  $\in$  STRING,
      delete ::= delete  $\in$  BOOLEAN ]],
    rollback ::= index  $\in$  Nat,
    targets ::= targets  $\in$  SUBSET (DOMAIN Target)
    phase     ::= phase  $\in$  Phase]
    status ::= status  $\in$  Status]

TYPE Proposal  $\triangleq$  [
  type      ::= type  $\in$  Type,
  index     ::= index  $\in$  Nat,
  values     ::= [ path  $\in$  SUBSET (DOMAIN Target[target].values)  $\mapsto$  [
    value ::= value  $\in$  STRING,
    delete ::= delete  $\in$  BOOLEAN ]],
    rollback ::= index  $\in$  Nat,
    dependencyIndex ::= dependencyIndex  $\in$  Nat,
    rollbackIndex ::= rollbackIndex  $\in$  Nat,
    rollbackValues ::= [ path  $\in$  SUBSET (DOMAIN Target[target].values)  $\mapsto$  [
      value ::= value  $\in$  STRING,
      delete ::= delete  $\in$  BOOLEAN ]],
    phase     ::= phase  $\in$  Phase]
    status     ::= status  $\in$  Status]

TYPE Configuration  $\triangleq$  [
  id        ::= id  $\in$  STRING,
  target    ::= target  $\in$  STRING,
  values     ::= [ path  $\in$  SUBSET (DOMAIN Target[target])  $\mapsto$  [
    value ::= value  $\in$  STRING,
    index ::= index  $\in$  Nat,
    deleted ::= delete  $\in$  BOOLEAN ]],
    configIndex ::= configIndex  $\in$  Nat,
    configTerm  ::= configTerm  $\in$  Nat,
    proposedIndex ::= proposedIndex  $\in$  Nat,
    committedIndex ::= committedIndex  $\in$  Nat,
    appliedIndex ::= appliedIndex  $\in$  Nat,
    appliedTerm  ::= appliedTerm  $\in$  Nat,

```

$$\begin{aligned} appliedValues ::= & [ path \in \text{SUBSET } (\text{DOMAIN } Target[target]) \mapsto [ \\ & value ::= value \in \text{STRING}, \\ & index ::= index \in \text{Nat}, \\ & deleted ::= delete \in \text{BOOLEAN} ]], \\ status ::= & status \in \text{Status} \end{aligned}$$

A transaction log. Transactions may either request a set of changes to a set of targets or rollback a prior change.

VARIABLE *transaction*

A record of per-target proposals

VARIABLE *proposal*

A record of per-target configurations

VARIABLE *configuration*

A record of target states

VARIABLE *target*

A record of target masterhips

VARIABLE *mastership*

$vars \triangleq \langle transaction, proposal, configuration, mastership, target \rangle$

This section models *mastership* for the configuration service.

Mastership is used primarily to track the lifecycle of individual configuration targets and react to state changes on the southbound. Each target is assigned a master from the *Node* set, and masters can be unset when the target disconnects.

Set node  $n$  as the master for target  $t$

$$\begin{aligned} SetMaster(n, t) &\triangleq \\ &\wedge mastership[t].master \neq n \\ &\wedge mastership' = [mastership \text{ EXCEPT } ![t].term = mastership[t].term + 1, \\ &\quad ![t].master = n] \\ &\wedge \text{UNCHANGED } \langle transaction, proposal, configuration, target \rangle \end{aligned}$$

$UnsetMaster(t) \triangleq$

$$\begin{aligned} &\wedge mastership[t].master \neq Nil \\ &\wedge mastership' = [mastership \text{ EXCEPT } ![t].master = Nil] \\ &\wedge \text{UNCHANGED } \langle transaction, proposal, configuration, target \rangle \end{aligned}$$

This section models configuration changes and rollbacks. Changes are appended to the transaction log and processed asynchronously.

$Value(s, t, p) \triangleq$

$$\begin{aligned} \text{LET } value &\triangleq \text{CHOOSE } v \in s : v.target = t \wedge v.path = p \\ \text{IN} \end{aligned}$$

$[value \mapsto value.value,$   
 $delete \mapsto value.delete]$

$Paths(s, t) \triangleq$   
 $[p \in \{v.path : v \in \{v \in s : v.target = t\}\} \mapsto Value(s, t, p)]$

$Changes(s) \triangleq$   
 $[t \in \{v.target : v \in s\} \mapsto Paths(s, t)]$

$ValidValues(t, p) \triangleq$   
 $UNION \{ \{ [value \mapsto v, delete \mapsto FALSE] : v \in Target[t].values[p] \}, \{ [value \mapsto Nil, delete \mapsto TRUE] \} \}$

$ValidPaths(t) \triangleq$   
 $UNION \{ \{ v @@@ [path \mapsto p] : v \in ValidValues(t, p) \} : p \in DOMAIN Target[t].values \}$

$ValidTargets \triangleq$   
 $UNION \{ \{ p @@@ [target \mapsto t] : p \in ValidPaths(t) \} : t \in DOMAIN Target \}$

The set of all valid sets of changes to all targets and their paths.

The set of possible changes is computed from the *Target* model value.

$ValidChanges \triangleq$   
 $LET changeSets \triangleq \{ s \in SUBSET ValidTargets :$   
 $\quad \forall t \in DOMAIN Target :$   
 $\quad \forall p \in DOMAIN Target[t].values :$   
 $\quad \quad Cardinality(\{ v \in s : v.target = t \wedge v.path = p \}) \leq 1 \}$   
 $IN$   
 $\{ Changes(s) : s \in changeSets \}$

The next available index in the transaction log.

This is computed as the max of the existing indexes in the log to  
allow for changes to the log (*e.g.* log compaction) to be modeled.

$NextIndex \triangleq$   
 $IF DOMAIN transaction = \{ \} THEN$   
 $1$   
 $ELSE$   
 $LET i \triangleq CHOOSE i \in DOMAIN transaction :$   
 $\quad \forall j \in DOMAIN transaction : i \geq j$   
 $IN i + 1$

Add a set of changes 'c' to the transaction log

$RequestChange(c) \triangleq$   
 $\wedge \exists isolation \in \{ ReadCommitted, Serializable \} :$   
 $\quad \wedge transaction' = transaction @@@ (NextIndex :> [type \mapsto Change,$   
 $\quad \quad \quad index \mapsto NextIndex,$   
 $\quad \quad \quad isolation \mapsto isolation,$   
 $\quad \quad \quad values \mapsto c,$   
 $\quad \quad \quad targets \mapsto \{ \},$

$$\begin{aligned}
& \text{phase} \mapsto \text{Initialize}, \\
& \text{status} \mapsto \text{Pending}] \\
& \wedge \text{UNCHANGED } \langle \text{proposal}, \text{configuration}, \text{mastership}, \text{target} \rangle \\
& \text{Add a rollback of transaction 't' to the transaction log} \\
& \text{RequestRollback}(t) \triangleq \\
& \wedge \exists \text{isolation} \in \{\text{ReadCommitted}, \text{Serializable}\} : \\
& \quad \wedge \text{transaction}' = \text{transaction} @@ (\text{NextIndex} :> [\text{type} \mapsto \text{Rollback}, \\
& \quad \text{index} \mapsto \text{NextIndex}, \\
& \quad \text{isolation} \mapsto \text{isolation}, \\
& \quad \text{rollback} \mapsto t, \\
& \quad \text{targets} \mapsto \{\}, \\
& \quad \text{phase} \mapsto \text{Initialize}, \\
& \quad \text{status} \mapsto \text{Pending}]) \\
& \wedge \text{UNCHANGED } \langle \text{proposal}, \text{configuration}, \text{mastership}, \text{target} \rangle
\end{aligned}$$


---

This section models the Transaction log reconciler.

Transactions come in two flavors : – *Change* transactions contain a set of changes to be applied to a set of *targets* – *Rollback* transactions reference a prior change transaction to be reverted to the previous state

Transactions proceed through a series of phases:

- \* *Initialize* – create and link Proposals
- \* *Validate* – validate changes and rollbacks
- \* *Commit* – commit changes to Configurations
- \* *Apply* – commit changes to Targets

Reconcile a transaction

$\text{ReconcileTransaction}(n, i) \triangleq$

Initialize is the only transaction phase that's globally serialized.

While in the *Initializing* phase, the reconciler checks whether the prior transaction has been *Initialized* before creating Proposals in the *Initialize* phase. Once all of the transaction's proposals have been *Initialized*, the transaction will be marked *Initialized*. If any proposal is *Failed*, the transaction will be marked *Failed* as well.

$\wedge \vee \wedge \text{transaction}[i].\text{phase} = \text{Initialize}$   
 $\wedge \vee \wedge \text{transaction}[i].\text{status} = \text{Pending}$

Serialize transaction initialization

$\wedge \vee i - 1 \notin \text{DOMAIN } \text{transaction}$   
 $\vee \text{Phase}[\text{transaction}[i - 1].\text{phase}] > \text{Phase}[\text{Initialize}]$   
 $\vee \text{transaction}[i - 1].\text{status} \neq \text{Pending}$

If the transaction's targets are not yet set, create proposals and add targets to the transaction state.

$\wedge \vee \wedge \text{transaction}[i].\text{targets} = \{\}$

If the transaction is a change, the targets are taken from the change values.

$\wedge \vee \wedge transaction[i].type = Change$   
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].targets = \text{DOMAIN } transaction[i].values]$   
 $\wedge proposal' = [t \in \text{DOMAIN } proposal \mapsto$   
 $\quad \text{IF } t \in \text{DOMAIN } transaction[i].values \text{ THEN}$   
 $\quad \quad proposal[t] @@ (i :> [type \mapsto Change,$   
 $\quad \quad \quad index \mapsto i,$   
 $\quad \quad \quad values \mapsto transaction[i].values[t],$   
 $\quad \quad \quad dependencyIndex \mapsto 0,$   
 $\quad \quad \quad rollbackIndex \mapsto 0,$   
 $\quad \quad \quad rollbackValues \mapsto \langle \rangle,$   
 $\quad \quad \quad phase \mapsto Initialize,$   
 $\quad \quad \quad status \mapsto Pending])$   
 $\quad \text{ELSE}$   
 $\quad \quad proposal[t]]$   
 If the transaction is a rollback, the targets affected are  
 the targets of the change transaction being rolled back.  
 $\vee \wedge transaction[i].type = Rollback$   
 $\wedge \vee \wedge transaction[i].rollback \in \text{DOMAIN } transaction$   
 $\wedge transaction[transaction[i].rollback].type = Change$   
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].targets =$   
 $\quad \text{DOMAIN } transaction[transaction[i].rollback].values]$   
 $\wedge proposal' = [t \in \text{DOMAIN } proposal \mapsto$   
 $\quad \text{IF } t \in \text{DOMAIN } transaction[transaction[i].rollback].values \text{ THEN}$   
 $\quad \quad proposal[t] @@ (i :> [type \mapsto Rollback,$   
 $\quad \quad \quad index \mapsto i,$   
 $\quad \quad \quad rollback \mapsto transaction[i].rollback,$   
 $\quad \quad \quad dependencyIndex \mapsto 0,$   
 $\quad \quad \quad rollbackIndex \mapsto 0,$   
 $\quad \quad \quad rollbackValues \mapsto \langle \rangle,$   
 $\quad \quad \quad phase \mapsto Initialize,$   
 $\quad \quad \quad status \mapsto Pending])$   
 $\quad \text{ELSE}$   
 $\quad \quad proposal[t]]$   
 $\vee \wedge \vee \wedge transaction[i].rollback \in \text{DOMAIN } transaction$   
 $\wedge transaction[transaction[i].rollback].type = Rollback$   
 $\vee transaction[i].rollback \notin \text{DOMAIN } transaction$   
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].status = Failed]$   
 $\wedge \text{UNCHANGED } \langle proposal \rangle$   
 $\vee \wedge transaction[i].targets \neq \{\}$   
 If all proposals have been *Complete*, mark the transaction *Complete*.  
 $\wedge \vee \wedge \forall t \in transaction[i].targets : proposal[t][i].status = Complete$   
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].status = Complete]$   
 $\wedge \text{UNCHANGED } \langle proposal \rangle$   
 If any proposal has been *Failed*, mark the transaction *Failed*.  
 $\vee \wedge \exists t \in transaction[i].targets : proposal[t][i].status = Failed$

$\wedge \text{transaction}' = [\text{transaction EXCEPT } ![i].\text{status} = \text{Failed}]$   
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$

Once the transaction has been *Initialized*, proceed to the *Validate* phase.  
If any of the transaction's proposals depend on a *Serializable* transaction,  
verify the dependency has been *Validated* to preserve serializability before  
moving the transaction to the *Validate* phase.

$\vee \wedge \text{transaction}[i].\text{status} = \text{Complete}$   
 $\wedge \forall t \in \text{transaction}[i].\text{targets} :$   
     $\vee \text{proposal}[t][i].\text{dependencyIndex} = 0$   
     $\vee \text{transaction}[\text{proposal}[t][i].\text{dependencyIndex}].\text{isolation} \neq \text{Serializable}$   
     $\vee \text{Phase}[\text{transaction}[\text{proposal}[t][i].\text{dependencyIndex}].\text{phase}] > \text{Phase}[\text{Validate}]$   
     $\vee \wedge \text{transaction}[\text{proposal}[t][i].\text{dependencyIndex}].\text{phase} = \text{Validate}$   
         $\wedge \text{transaction}[\text{proposal}[t][i].\text{dependencyIndex}].\text{status} \in \{\text{Complete}, \text{Failed}\}$   
 $\wedge \text{transaction}' = [\text{transaction EXCEPT } ![i].\text{phase} = \text{Validate},$   
 $![i].\text{status} = \text{Pending}]$   
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$

$\vee \wedge \text{transaction}[i].\text{status} = \text{Failed}$   
 $\wedge \text{transaction}' = [\text{transaction EXCEPT } ![i].\text{phase} = \text{Abort},$   
 $![i].\text{status} = \text{Pending}]$   
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$

$\vee \wedge \text{transaction}[i].\text{phase} = \text{Validate}$   
 $\wedge \vee \wedge \text{transaction}[i].\text{status} = \text{Pending}$   
    Move the transaction's proposals to the *Validating* state  
 $\wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : \text{proposal}[t][i].\text{phase} \neq \text{Validate}$   
     $\wedge \text{proposal}' = [t \in \text{DOMAIN proposal} \mapsto$   
 $\text{IF } t \in \text{transaction}[i].\text{targets} \text{ THEN}$   
 $[\text{proposal}[t] \text{ EXCEPT } ![i].\text{phase} = \text{Validate},$   
 $![i].\text{status} = \text{Pending}]$   
 $\text{ELSE}$   
 $\text{proposal}[t]]$   
     $\wedge \text{UNCHANGED } \langle \text{transaction} \rangle$   
    If all proposals have been *Complete*, mark the transaction *Complete*.  
 $\vee \wedge \forall t \in \text{transaction}[i].\text{targets} : \text{proposal}[t][i].\text{status} = \text{Complete}$   
     $\wedge \text{transaction}' = [\text{transaction EXCEPT } ![i].\text{status} = \text{Complete}]$   
     $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$   
    If any proposal has been *Failed*, mark the transaction *Failed*.  
 $\vee \wedge \exists t \in \text{transaction}[i].\text{targets} : \text{proposal}[t][i].\text{status} = \text{Failed}$   
     $\wedge \text{transaction}' = [\text{transaction EXCEPT } ![i].\text{status} = \text{Failed}]$   
     $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$

Once the transaction has been *Validated*, proceed to the *Commit* phase.  
If any of the transaction's proposals depend on a *Serializable* transaction,  
verify the dependency has been *Committed* to preserve serializability before  
moving the transaction to the *Commit* phase.

$\vee \wedge \text{transaction}[i].\text{status} = \text{Complete}$   
 $\wedge \forall t \in \text{transaction}[i].\text{targets} :$

$$\vee \wedge transaction[i].status = Complete$$
$$\forall \text{ proposal}[t][i]. \text{dependencyIndex} = 0$$
$$\forall transaction[proposal[t][i].dependencyIndex].isolation \neq Serializable$$
$$\vee Phase[transaction[proposal[t]][i].dependencyIndex].phase] > Phase[proposal[t][i].dependencyIndex].phase]$$
$$\wedge transaction[proposal[t][i].dependencyIndex].phase = Validate$$
$$\wedge transaction[proposal[t][i].dependencyIndex].status \in \{Complete, Failed\}$$
$$\begin{aligned} & ! [i].phase = \textit{Validate}, \\ & ! [i].status = \textit{Pending} \end{aligned}$$
 $\wedge$  UNCHANGED  $\langle proposal \rangle$ 
$$\vee \wedge transaction[i].status = Failed$$
$$\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Abort,$$
$$! [i].status = Pending]$$
 $\wedge$  UNCHANGED  $\langle proposal \rangle$ 
$$\vee \wedge transaction[i].phase = Validate$$
$$\wedge \vee \wedge transaction[i].status = Pending$$

Move the transaction's proposals to the *Validating* state

$$\wedge \vee \wedge \exists t \in transaction[i].targets : proposal[t][i].phase \neq Validate$$
$$\wedge proposal' = [t \in \text{DOMAIN } proposal \mapsto$$

IF  $t \in transaction[i].targets$  THEN

$$[proposal[t] \text{ EXCEPT } ![i].phase = \text{Validate},$$
$$! [i].status = Pending]$$

ELSE

$$proposal[t]]$$
 $\wedge$  UNCHANGED  $\langle transaction \rangle$ 

If all proposals have been *Complete*, mark the transaction *Complete*.

$$\forall i \wedge \forall t \in transaction[i].targets : proposal[t][i].status = Complete$$
$$\wedge transaction' = [transaction \text{ EXCEPT } ! [i].status = Complete]$$
 $\wedge$  UNCHANGED  $\langle proposal \rangle$ 

If any proposal has been *Failed*, mark the transaction *Failed*.

$$\forall \wedge \exists t \in transaction[i].targets : proposal[t][i].status = Failed$$
$$\wedge transaction' = [transaction \text{ EXCEPT } ! [i].status = Failed]$$
 $\wedge$  UNCHANGED  $\langle proposal \rangle$ 

Once the transaction has been *Validated*, proceed to the *Commit* phase.

If any of the transaction's proposals depend on a *Serializable* transaction,

verify the dependency has been *Committed* to preserve serializability before

moving the transaction to the *Commit* phase.

$$\vee \wedge transaction[i].status = Complete$$
$$\wedge \forall t \in transaction[i].targets :$$





ELSE

$proposal[t]$

$\wedge$  UNCHANGED  $\langle transaction \rangle$

If all proposals have been *Complete*, mark the transaction *Complete*.

$\vee \wedge \forall t \in transaction[i].targets : proposal[t][i].status = Complete$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].status = Complete]$

$\wedge$  UNCHANGED  $\langle proposal \rangle$

If any proposal has been *Failed*, mark the transaction *Failed*.

$\vee \wedge \exists t \in transaction[i].targets : proposal[t][i].status = Failed$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].status = Failed]$

$\wedge$  UNCHANGED  $\langle proposal \rangle$

The *Aborting* state is used to clean up transactions that have failed during the *Initializing* or *Validating* phases.

$\vee \wedge transaction[i].phase = Abort$

$\wedge transaction[i].status = Pending$

Move the transaction's proposals to the *Aborting* state

$\wedge \vee \wedge \exists t \in transaction[i].targets : proposal[t][i].phase \neq Abort$

$\wedge proposal' = [t \in \text{DOMAIN } proposal \mapsto$

IF  $t \in transaction[i].targets$  THEN

$[proposal[t] \text{ EXCEPT } ![i].phase = Abort,$

$![i].status = Pending]$

ELSE

$proposal[t]$

$\wedge$  UNCHANGED  $\langle transaction \rangle$

If all proposals have been *Complete*, mark the transaction *Complete*.

$\vee \wedge \forall t \in transaction[i].targets : proposal[t][i].status = Complete$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].status = Complete]$

$\wedge$  UNCHANGED  $\langle proposal \rangle$

$\wedge$  UNCHANGED  $\langle configuration, mastership, target \rangle$

Reconcile a proposal

$ReconcileProposal(n, t, i) \triangleq$

$\wedge \vee \wedge proposal[t][i].phase = Initialize$

$\wedge proposal[t][i].status = Pending$

$\wedge proposal' = [proposal \text{ EXCEPT } ![t] = [$

$proposal[t] \text{ EXCEPT } ![i] = [$

$status \mapsto Complete,$

$dependencyIndex \mapsto configuration[t].proposedIndex] @@ proposal[t][i]]$

$\wedge configuration' = [configuration \text{ EXCEPT } ![t].proposedIndex = i]$

$\wedge$  UNCHANGED  $\langle target \rangle$

While in the *Validate* phase, validate the proposed changes.

If validation is successful, the proposal also records the changes required to roll back the proposal and the index to which to roll back.

$\vee \wedge proposal[t][i].phase = Validate$

$\wedge proposal[t][i].status = Pending$

$\wedge \text{configuration}[t].\text{committedIndex} = \text{proposal}[t][i].\text{dependencyIndex}$   
 For *Change* proposals validate the set of requested changes.  
 $\wedge \vee \wedge \text{proposal}[t][i].\text{type} = \text{Change}$   
 $\wedge \text{LET } \text{rollbackIndex} \triangleq \text{configuration}[t].\text{configIndex}$   
 $\text{rollbackValues} \triangleq [p \in \text{DOMAIN } \text{proposal}[t][i].\text{values} \mapsto$   
     IF  $p \in \text{DOMAIN } \text{configuration}[t].\text{values}$  THEN  
          $\text{configuration}[t].\text{values}[p]$   
     ELSE  
          $[value \mapsto \text{Nil},$   
          $\text{delete} \mapsto \text{TRUE}]$   
 Model validation successes and failures with *Valid* and *Invalid* results.  
 IN  $\exists r \in \{ \text{Valid}, \text{Invalid} \} :$   
     If the *Change* is *Valid*, record the changes required to roll  
     back the proposal and the index to which the rollback changes  
     will roll back the configuration.  
 $\vee \wedge r = \text{Valid}$   
 $\wedge \text{proposal}' = [\text{proposal EXCEPT } ![t] = [$   
      $\text{proposal}[t] \text{ EXCEPT } ![i].\text{rollbackIndex} = \text{rollbackIndex},$   
      $![i].\text{rollbackValues} = \text{rollbackValues},$   
      $![i].\text{status} = \text{Complete}]$   
 $\vee \wedge r = \text{Invalid}$   
 $\wedge \text{proposal}' = [\text{proposal EXCEPT } ![t] = [$   
      $\text{proposal}[t] \text{ EXCEPT } ![i].\text{status} = \text{Failed}]$   
 For *Rollback* proposals, validate the rollback changes which are  
 proposal being rolled back.  
 $\vee \wedge \text{proposal}[t][i].\text{type} = \text{Rollback}$   
 Rollbacks can only be performed on *Change* type proposals.  
 $\wedge \vee \wedge \text{proposal}[t][\text{proposal}[t][i].\text{rollback}].\text{type} = \text{Change}$   
 Only roll back the change if it's the latest change made  
 to the configuration based on the configuration index.  
 $\wedge \vee \wedge \text{configuration}[t].\text{configIndex} = \text{proposal}[t][i].\text{rollback}$   
 $\wedge \text{LET } \text{rollbackIndex} \triangleq \text{proposal}[t][\text{proposal}[t][i].\text{rollback}].\text{rollbackIndex}$   
 $\text{rollbackValues} \triangleq \text{proposal}[t][\text{proposal}[t][i].\text{rollback}].\text{rollbackValues}$   
 IN  $\exists r \in \{ \text{Valid}, \text{Invalid} \} :$   
     If the *Rollback* is *Valid*, record the changes required to  
     roll back the target proposal and the index to which the  
     configuration is being rolled back.  
 $\vee \wedge r = \text{Valid}$   
 $\wedge \text{proposal}' = [\text{proposal EXCEPT } ![t] = [$   
      $\text{proposal}[t] \text{ EXCEPT } ![i].\text{rollbackIndex} = \text{rollbackIndex},$   
      $![i].\text{rollbackValues} = \text{rollbackValues},$   
      $![i].\text{status} = \text{Complete}]$   
 $\vee \wedge r = \text{Invalid}$   
 $\wedge \text{proposal}' = [\text{proposal EXCEPT } ![t] = [$   
      $\text{proposal}[t] \text{ EXCEPT } ![i].\text{status} = \text{Failed}]$

If the *Rollback* target is not the most recent change to the configuration,  
fail validation for the proposal.

$$\vee \wedge \text{configuration}[t].\text{configIndex} \neq \text{proposal}[t][i].\text{rollback}$$

$$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = [\text{proposal}[t] \text{ EXCEPT } ![i].\text{status} = \text{Failed}]]$$

If a *Rollback* proposal is attempting to roll back another *Rollback*,  
fail validation for the proposal.

$$\vee \wedge \text{proposal}[t][\text{proposal}[t][i].\text{rollback}].\text{type} = \text{Rollback}$$

$$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = [$$

$$\quad \text{proposal}[t] \text{ EXCEPT } ![i].\text{status} = \text{Failed}]]$$

$$\wedge \text{UNCHANGED } \langle \text{configuration}, \text{target} \rangle$$

While in the *Commit* state, commit the proposed changes to the configuration.

$$\vee \wedge \text{proposal}[t][i].\text{phase} = \text{Commit}$$

$$\wedge \text{proposal}[t][i].\text{status} = \text{Pending}$$

Only commit the proposal if the prior proposal has already been committed.

$$\wedge \text{configuration}[t].\text{committedIndex} = \text{proposal}[t][i].\text{dependencyIndex}$$

If the proposal is a change, commit the change values and set the configuration  
index to the proposal index.

$$\wedge \vee \wedge \text{proposal}[t][i].\text{type} = \text{Change}$$

$$\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{values} = \text{proposal}[t][i].\text{values},$$

$$\quad \quad \quad ![t].\text{configIndex} = i,$$

$$\quad \quad \quad ![t].\text{committedIndex} = i]$$

If the proposal is a rollback, commit the rollback values and index. This  
will cause the configuration index to be reverted to the index prior to  
the transaction/proposal being rolled back.

$$\vee \wedge \text{proposal}[t][i].\text{type} = \text{Rollback}$$

$$\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{values} = \text{proposal}[t][i].\text{rollbackValues},$$

$$\quad \quad \quad ![t].\text{configIndex} = \text{proposal}[t][i].\text{rollbackIndex},$$

$$\quad \quad \quad ![t].\text{committedIndex} = i]$$

$$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = [\text{proposal}[t] \text{ EXCEPT } ![i].\text{status} = \text{Complete}]]$$

$$\wedge \text{UNCHANGED } \langle \text{target} \rangle$$

While in the *Apply* phase, apply the proposed changes to the target.

$$\vee \wedge \text{proposal}[t][i].\text{phase} = \text{Apply}$$

$$\wedge \text{proposal}[t][i].\text{status} = \text{Pending}$$

$$\wedge \text{configuration}[t].\text{appliedIndex} = \text{proposal}[t][i].\text{dependencyIndex}$$

$$\wedge \text{configuration}[t].\text{appliedTerm} = \text{mastership}[t].\text{term}$$

$$\wedge \text{mastership}[t].\text{master} = n$$

Model successful and failed target update requests.

$$\wedge \exists r \in \{\text{Success}, \text{Failure}\} :$$

$$\vee \wedge r = \text{Success}$$

If the proposal is a change, apply the change values to the target  
and update the configuration's applied index and values.

$$\wedge \vee \wedge \text{proposal}[t][i].\text{type} = \text{Change}$$

$$\wedge \text{target}' = [\text{target} \text{ EXCEPT } ![t] = \text{proposal}[t][i].\text{values} @@ \text{target}[t]]$$

$$\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT }$$

$$\quad \quad \quad ![t].\text{appliedIndex} = i,$$



Formal specification, constraints, and theorems.

$$\begin{aligned}
Next &\triangleq \\
&\vee \exists c \in ValidChanges : \\
&\quad RequestChange(c) \\
&\vee \exists t \in DOMAIN\ trans : \\
&\quad RequestRollback(t) \\
&\vee \exists n \in Node : \\
&\quad \exists t \in DOMAIN\ Target : \\
&\quad \quad SetMaster(n, t) \\
&\quad \vee \exists t \in DOMAIN\ Target : \\
&\quad \quad UnsetMaster(t) \\
&\vee \exists n \in Node : \\
&\quad \exists t \in DOMAIN\ trans : \\
&\quad \quad ReconcileTranso
\end{aligned}$$

$$\begin{aligned}
& \forall \exists n \in \text{Node} : \\
& \quad \exists t \in \text{DOMAIN } \text{proposal} : \\
& \quad \quad \exists i \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \quad \text{ReconcileProposal}(n, t, i) \\
& \forall \exists n \in \text{Node} : \\
& \quad \exists c \in \text{DOMAIN } \text{configuration} : \\
& \quad \quad \text{ReconcileConfiguration}(n, c) \\
\text{Spec} & \triangleq \text{Init} \wedge \Box[\text{Next}]_{\text{vars}} \\
\text{Order} & \triangleq \\
& \wedge \forall i, j \in \text{DOMAIN } \text{transaction} : \\
& \quad \vee j \leq i \\
& \quad \vee \text{Phase}[\text{transaction}[i].\text{phase}] \geq \text{Phase}[\text{transaction}[j].\text{phase}] \\
& \quad \vee \text{transaction}[j].\text{status} = \text{Failed} \\
& \wedge \forall t \in \text{DOMAIN } \text{proposal} : \\
& \quad \forall i, j \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \vee j \leq i \\
& \quad \quad \vee \text{Phase}[\text{proposal}[t][i].\text{phase}] \geq \text{Phase}[\text{proposal}[t][j].\text{phase}] \\
& \quad \quad \vee \text{proposal}[t][i].\text{status} = \text{Failed} \\
\text{Consistency} & \triangleq \\
& \forall t \in \text{DOMAIN } \text{target} : \\
& \quad \text{LET} \\
& \quad \quad \text{Compute the transaction indexes that have been applied to the target} \\
& \quad \quad \text{appliedIndexes} \triangleq \{i \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \quad \wedge \text{transaction}[i].\text{type} = \text{Change} \\
& \quad \quad \quad \wedge i \in \text{DOMAIN } \text{proposal}[t] \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{phase} = \text{Apply} \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{status} = \text{Complete} \\
& \quad \quad \quad \wedge t \in \text{DOMAIN } \text{transaction}[i].\text{values} \\
& \quad \quad \quad \wedge \neg \exists j \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \quad \quad \wedge j > i \\
& \quad \quad \quad \quad \wedge \text{transaction}[j].\text{type} = \text{Rollback} \\
& \quad \quad \quad \quad \wedge \text{transaction}[j].\text{rollback} = i \\
& \quad \quad \quad \quad \wedge \text{transaction}[j].\text{phase} = \text{Apply} \\
& \quad \quad \quad \quad \wedge \text{transaction}[j].\text{status} = \text{Complete}\} \\
& \quad \quad \text{Compute the set of paths in the target that have been updated by transactions} \\
& \quad \quad \text{appliedPaths} \triangleq \text{UNION } \{\text{DOMAIN } \text{transaction}[i].\text{values}[t] : i \in \text{appliedIndexes}\} \\
& \quad \quad \text{Compute the highest index applied to the target for each path} \\
& \quad \quad \text{pathIndexes} \triangleq [p \in \text{appliedPaths} \mapsto \text{CHOOSE } i \in \text{appliedIndexes} : \\
& \quad \quad \quad \forall j \in \text{appliedIndexes} : \\
& \quad \quad \quad \quad \wedge i \geq j \\
& \quad \quad \quad \quad \wedge p \in \text{DOMAIN } \text{transaction}[i].\text{values}] \\
& \quad \quad \text{Compute the expected target configuration based on the last indexes applied} \\
& \quad \quad \text{to the target for each path.}
\end{aligned}$$

$$\begin{aligned}
& \text{expectedConfig} \triangleq [p \in \text{DOMAIN } \text{pathIndexes} \mapsto \text{transaction}[\text{pathIndexes}[p]].\text{values}[p]] \\
& \text{IN} \\
& \text{target}[t] = \text{expectedConfig} \\
& \text{Isolation} \triangleq \\
& \quad \forall i, j \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \vee j \leq i \\
& \quad \quad \vee \text{transaction}[i].\text{targets} \cap \text{transaction}[j].\text{targets} = \{\} \\
& \quad \quad \vee \text{transaction}[i].\text{isolation} \neq \text{Serializable} \\
& \quad \quad \vee \wedge \vee \wedge \text{transaction}[i].\text{phase} \in \{\text{Commit}, \text{Abort}\} \\
& \quad \quad \quad \wedge \text{transaction}[i].\text{status} \in \{\text{Complete}, \text{Failed}\} \\
& \quad \quad \quad \vee \text{Phase}[\text{transaction}[i].\text{phase}] > \text{Phase}[\text{Commit}] \\
& \quad \quad \quad \vee \text{Phase}[\text{transaction}[j].\text{phase}] < \text{Phase}[\text{Commit}] \\
& \quad \quad \wedge \vee \wedge \text{transaction}[i].\text{phase} \in \{\text{Apply}, \text{Abort}\} \\
& \quad \quad \quad \wedge \text{transaction}[i].\text{status} \in \{\text{Complete}, \text{Failed}\} \\
& \quad \quad \quad \vee \text{Phase}[\text{transaction}[j].\text{phase}] < \text{Phase}[\text{Apply}] \\
& \quad \quad \vee \text{transaction}[j].\text{status} = \text{Failed} \\
& \text{THEOREM } \text{Safety} \triangleq \text{Spec} \Rightarrow \Box(\text{Order} \wedge \text{Consistency} \wedge \text{Isolation}) \\
& \text{Completion} \triangleq \\
& \quad \wedge \forall i \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \wedge \text{transaction}[i].\text{phase} = \text{Commit} \\
& \quad \quad \wedge \text{transaction}[i].\text{status} \in \{\text{Complete}, \text{Failed}\} \\
& \quad \wedge \forall i \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \wedge \text{transaction}[i].\text{phase} = \text{Apply} \\
& \quad \quad \wedge \text{transaction}[i].\text{status} \in \{\text{Complete}, \text{Failed}\} \\
& \quad \wedge \forall t \in \text{DOMAIN } \text{proposal} : \\
& \quad \quad \forall i \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{phase} = \text{Commit} \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{status} \in \{\text{Complete}, \text{Failed}\} \\
& \quad \wedge \forall t \in \text{DOMAIN } \text{proposal} : \\
& \quad \quad \forall i \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{phase} = \text{Apply} \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{status} \in \{\text{Complete}, \text{Failed}\}
\end{aligned}$$

THEOREM *Liveness*  $\triangleq \text{Spec} \Rightarrow \Diamond \text{Completion}$

---

Type assumptions.

ASSUME  $\text{Nil} \in \text{STRING}$

ASSUME  $\forall \text{phase} \in \text{Phase} : \text{phase} \in \text{STRING}$

ASSUME  $\forall \text{status} \in \text{Status} : \text{status} \in \text{STRING}$

ASSUME  $\wedge \text{IsFiniteSet}(\text{Node})$



```

       $\wedge \forall n \in Node :$ 
         $\wedge n \notin \text{DOMAIN } Target$ 
         $\wedge n \in \text{STRING}$ 

ASSUME  $\wedge \forall t \in \text{DOMAIN } Target :$ 
       $\wedge t \notin Node$ 
       $\wedge t \in \text{STRING}$ 
       $\wedge Target[t].persistent \in \text{BOOLEAN}$ 
       $\wedge \forall p \in \text{DOMAIN } Target[t].values :$ 
         $IsFiniteSet(Target[t].values[p])$ 

```

---

```

\ * Modification History
\ * Last modified Mon Feb 07 02:17:42 PST 2022 by jordanhaltermann
\ * Created Wed Sep 22 13:22:32 PDT 2021 by jordanhaltermann

```