
MODULE *Config*

INSTANCE *Naturals*

INSTANCE *FiniteSets*

INSTANCE *Sequences*

INSTANCE *TLC*

Indicates that a configuration change is waiting to be applied to the network
CONSTANT *Pending*

Indicates that a configuration change has been applied to the network
CONSTANT *Complete*

Indicates that a configuration change failed
CONSTANT *Failed*

Indicates a change is a configuration
CONSTANT *Change*

Indicates a change is a rollback
CONSTANT *Rollback*

Indicates a device is connected
CONSTANT *Connected*

Indicates a device is disconnected
CONSTANT *Disconnected*

Indicates that an error occurred when applying a change
CONSTANT *Error*

An empty constant
CONSTANT *Nil*

The set of all nodes
CONSTANT *Node*

The set of all devices
CONSTANT *Device*

ASSUME *Pending* ∈ STRING

ASSUME *Complete* ∈ STRING

ASSUME *Failed* ∈ STRING

ASSUME *Rollback* ∈ STRING

ASSUME *Connected* ∈ STRING

ASSUME $Disconnected \in \text{STRING}$
 ASSUME $Error \in \text{STRING}$
 ASSUME $Nil \in \text{STRING}$

ASSUME $\wedge IsFiniteSet(Node)$
 $\wedge \forall n \in Node :$
 $\wedge n \notin Device$
 $\wedge n \in \text{STRING}$
 ASSUME $\wedge IsFiniteSet(Device)$
 $\wedge \forall d \in Device :$
 $\wedge d \notin Node$
 $\wedge d \in \text{STRING}$

Per-node election state
 VARIABLE $leader$

Per-node per-device election state
 VARIABLE $master$

A sequence of network-wide configuration changes
 Each change contains a record of 'changes' for each device
 VARIABLE $networkChange$

A record of sequences of device configuration changes
 Each sequence is a list of changes in the order in which they
 are to be applied to the device
 VARIABLE $deviceChange$

A record of device states - either Available or Unavailable
 VARIABLE $deviceState$

A count of leader changes to serve as a state constraint
 VARIABLE $electionCount$

A count of configuration changes to serve as a state constraint
 VARIABLE $configCount$

A count of device connection changes to serve as a state constraint
 VARIABLE $connectionCount$

Node variables
 $nodeVars \triangleq \langle leader, master \rangle$

Configuration variables
 $configVars \triangleq \langle networkChange, deviceChange \rangle$

Device variables
 $deviceVars \triangleq \langle deviceState \rangle$

State constraint variables
 $constraintVars \triangleq \langle electionCount, configCount, connectionCount \rangle$

$vars \triangleq \langle nodeVars, configVars, deviceVars, constraintVars \rangle$

This section models leader election for control loops and for devices. Leader election is modelled as a simple boolean indicating whether each node is the leader for the cluster and for each device. This model implies the ordering of leadership changes is irrelevant to the correctness of the spec.

Set the leader for node n to l
 $SetNodeLeader(n, l) \triangleq$
 $\wedge leader' = [leader \text{ EXCEPT } ![n] = n = l]$
 $\wedge electionCount' = electionCount + 1$
 $\wedge \text{UNCHANGED } \langle master, configVars, deviceVars, configCount, connectionCount \rangle$

Set the master for device d on node n to l
 $SetDeviceMaster(n, d, l) \triangleq$
 $\wedge master' = [master \text{ EXCEPT } ![n] = [master[n] \text{ EXCEPT } ![d] = n = l]]$
 $\wedge electionCount' = electionCount + 1$
 $\wedge \text{UNCHANGED } \langle leader, configVars, deviceVars, configCount, connectionCount \rangle$

This section models the northbound API for the configuration service.

Enqueue network configuration change c
 $SubmitChange(c) \triangleq$
 $\wedge Cardinality(\text{DOMAIN } c) > 0$
 $\wedge networkChange' = Append(networkChange, [$
 $\quad phase \quad \mapsto Change,$
 $\quad changes \quad \mapsto c,$
 $\quad value \quad \mapsto Len(networkChange),$
 $\quad state \quad \mapsto Pending,$
 $\quad incarnation \mapsto 0])$
 $\wedge configCount' = configCount + 1$
 $\wedge \text{UNCHANGED } \langle nodeVars, deviceChange, deviceVars, electionCount, connectionCount \rangle$

$RollbackChange(c) \triangleq$
 $\wedge networkChange[c].phase = Change$
 $\wedge networkChange[c].state = Complete$
 $\wedge networkChange' = [networkChange \text{ EXCEPT } ![c].phase = Rollback, ![c].state = Pending]$
 $\wedge configCount' = configCount + 1$
 $\wedge \text{UNCHANGED } \langle nodeVars, deviceChange, deviceVars, electionCount, connectionCount \rangle$

This section models the *NetworkChange* reconciler. The reconciler reconciles network changes when the change or one of its device changes is updated.

Return the set of all network changes prior to the given change

$$\begin{aligned} \text{PriorNetworkChanges}(c) &\triangleq \\ &\{n \in \text{DOMAIN } \text{networkChange} : n < c\} \end{aligned}$$

Return the set of all completed device changes for network change c

$$\begin{aligned} \text{NetworkCompletedChanges}(c) &\triangleq \\ &\{d \in \text{DOMAIN } \text{networkChange}[c].\text{changes} : \\ &\quad \wedge c \in \text{DOMAIN } \text{deviceChange}[d] \\ &\quad \wedge \text{deviceChange}[d][c].\text{state} = \text{Complete}\} \end{aligned}$$

Return a boolean indicating whether all device changes are complete for the given network change

$$\begin{aligned} \text{NetworkChangesComplete}(c) &\triangleq \\ &\text{Cardinality}(\text{NetworkCompletedChanges}(c)) = \text{Cardinality}(\text{DOMAIN } \text{networkChange}[c].\text{changes}) \end{aligned}$$

Return the set of all incomplete device changes prior to network change c

$$\begin{aligned} \text{PriorIncompleteDevices}(c) &\triangleq \\ &\text{UNION } \{ \text{DOMAIN } \text{networkChange}[n].\text{changes} : \\ &\quad n \in \{n \in \text{PriorNetworkChanges}(c) : \neg \text{NetworkChangesComplete}(n)\} \} \end{aligned}$$

Return the set of all devices configured by network change c

$$\text{NetworkChangeDevices}(c) \triangleq \text{DOMAIN } \text{networkChange}[c].\text{changes}$$

Return the set of all connected devices configured by network change c

$$\text{ConnectedDevices}(c) \triangleq \{d \in \text{DOMAIN } \text{networkChange}[c].\text{changes} : \text{deviceState}[d] = \text{Connected}\}$$

Return a boolean indicating whether network change c can be applied

A change can be applied if its devices do not intersect with past device changes that have not been applied

$$\begin{aligned} \text{CanApplyNetworkChange}(c) &\triangleq \\ &\wedge \text{Cardinality}(\text{ConnectedDevices}(c) \cap \text{NetworkChangeDevices}(c)) \neq 0 \\ &\wedge \text{Cardinality}(\text{NetworkChangeDevices}(c) \cap \text{PriorIncompleteDevices}(c)) = 0 \\ &\wedge \forall \text{networkChange}[c].\text{incarnation} = 0 \\ &\quad \vee \text{Cardinality}(\{d \in \text{DOMAIN } \text{networkChange}[c].\text{changes} : \\ &\quad \quad \wedge \text{deviceChange}[d][c].\text{incarnation} = \text{networkChange}[c].\text{incarnation} \\ &\quad \quad \wedge \text{deviceChange}[d][c].\text{phase} = \text{Rollback} \\ &\quad \quad \wedge \text{deviceChange}[d][c].\text{state} = \text{Complete}\}) = \\ &\quad \text{Cardinality}(\text{DOMAIN } \text{networkChange}[c].\text{changes}) \end{aligned}$$

Return a boolean indicating whether a change exists for the given device

If the device is modified by the change, it must contain a device change that's either *Complete* or with the same 'incarnation' as the network change.

$$\begin{aligned} \text{HasDeviceChange}(d, c) &\triangleq \\ &\wedge c \in \text{DOMAIN } \text{deviceChange}[d] \\ &\wedge \text{deviceChange}[d][c].\text{incarnation} = \text{networkChange}[c].\text{incarnation} \end{aligned}$$

Return a boolean indicating whether device changes have been propagated
for the given network change

$$HasDeviceChanges(c) \triangleq$$

$$Cardinality(\{d \in \text{DOMAIN } networkChange[c].changes : HasDeviceChange(d, c)\}) =$$

$$Cardinality(\text{DOMAIN } networkChange[c].changes)$$

Add or update the given device changes for the given network change.
If a device change already exists, update the 'incarnation' field.

$$CreateDeviceChange(d, c) \triangleq$$

IF $d \in \text{DOMAIN } networkChange[c].changes$ THEN

IF $c \in \text{DOMAIN } deviceChange[d]$ THEN

IF $deviceChange[d][c].state = Complete$ THEN

$$deviceChange[d]$$

ELSE

$$[deviceChange[d] \text{ EXCEPT } ![c].incarnation = networkChange[c].incarnation,$$

$$![c].state = Pending]$$

ELSE

$$[x \in \{c\} \mapsto [$$

$$phase \mapsto networkChange[c].phase,$$

$$state \mapsto Pending,$$

$$value \mapsto networkChange[c].value,$$

$$incarnation \mapsto networkChange[c].incarnation]] @@ deviceChange[d]$$

ELSE

$$deviceChange[d]$$

Add or update device changes for the given network change

$$CreateDeviceChanges(c) \triangleq$$

$$deviceChange' = [d \in \text{DOMAIN } deviceChange \mapsto CreateDeviceChange(d, c)]$$

Rollback device change c for device d

$$RollbackDeviceChange(d, c) \triangleq$$

IF $\wedge c \in \text{DOMAIN } deviceChange[d]$

$$\wedge \vee deviceChange[d][c].phase = Change$$

$$\vee \wedge deviceChange[d][c].phase = Rollback$$

$$\wedge deviceChange[d][c].state = Failed$$

THEN

$$[deviceChange[d] \text{ EXCEPT } ![c].phase = Rollback, ![c].state = Pending]$$

ELSE

$$deviceChange[d]$$

Roll back device changes

$$RollbackDeviceChanges(c) \triangleq$$

$$deviceChange' = [d \in \text{DOMAIN } deviceChange \mapsto RollbackDeviceChange(d, c)]$$

Return a boolean indicating whether the given device change is *Failed*

$$IsFailedDeviceChange(d, c) \triangleq$$

$\wedge c \in \text{DOMAIN } deviceChange[d]$
 $\wedge deviceChange[d][c].incarnation = networkChange[c].incarnation$
 $\wedge deviceChange[d][c].state = Failed$

Return a boolean indicating whether the given device change is *Complete*

$IsCompleteDeviceChange(d, c) \triangleq$
 $\wedge c \in \text{DOMAIN } deviceChange[d]$
 $\wedge deviceChange[d][c].incarnation = networkChange[c].incarnation$
 $\wedge deviceChange[d][c].phase = Change$
 $\wedge deviceChange[d][c].state = Complete$

Return a boolean indicating whether any device change is *Failed* for the given network change

$HasFailedDeviceChanges(c) \triangleq$
 $Cardinality(\{d \in \text{DOMAIN } networkChange[c].changes :$
 $IsFailedDeviceChange(d, c)\}) \neq 0$

Return a boolean indicating whether all device changes are *Complete* for the given network change

$DeviceChangesComplete(c) \triangleq$
 $Cardinality(\{d \in \text{DOMAIN } networkChange[c].changes :$
 $IsCompleteDeviceChange(d, c)\}) =$
 $Cardinality(\text{DOMAIN } networkChange[c].changes)$

Reconcile a network change state

$ReconcileNetworkChange(n, c) \triangleq$
 $\wedge leader[n]$
 $\wedge networkChange[c].state = Pending$
 $\wedge \vee \wedge \neg HasDeviceChanges(c)$
 $\wedge CreateDeviceChanges(c)$
 $\wedge \text{UNCHANGED } \langle networkChange \rangle$
 $\vee \wedge HasDeviceChanges(c)$
 $\wedge \vee \wedge networkChange[c].phase = Change$
 $\wedge \vee \wedge CanApplyNetworkChange(c)$
 $\wedge networkChange' = [networkChange \text{ EXCEPT}$
 $! [c].incarnation = networkChange[c].incarnation + 1]$
 $\wedge \text{UNCHANGED } \langle deviceChange \rangle$
 $\vee \wedge DeviceChangesComplete(c)$
 $\wedge networkChange' = [networkChange \text{ EXCEPT}$
 $! [c].state = Complete]$
 $\wedge \text{UNCHANGED } \langle deviceChange \rangle$
 $\vee \wedge HasFailedDeviceChanges(c)$
 $\wedge RollbackDeviceChanges(c)$
 $\wedge \text{UNCHANGED } \langle networkChange \rangle$
 $TODO$
 $\vee \wedge networkChange[c].phase = Rollback$
 $\wedge networkChange' = [networkChange \text{ EXCEPT}$
 $! [c].state = Complete]$

$$\begin{aligned} & \wedge \text{UNCHANGED } \langle \text{deviceChange} \rangle \\ & \wedge \text{UNCHANGED } \langle \text{nodeVars}, \text{deviceVars}, \text{constraintVars} \rangle \end{aligned}$$

This section models the *DeviceChange* reconciler.

$$\begin{aligned} \text{ReconcileDeviceChange}(n, d, c) & \triangleq \\ & \wedge \text{master}[n][d] \\ & \wedge \text{deviceChange}[d][c].\text{state} = \text{Pending} \\ & \wedge \text{deviceChange}[d][c].\text{incarnation} > 0 \\ & \wedge \vee \wedge \text{deviceState}[d] = \text{Connected} \\ & \quad \wedge \text{deviceChange}' = [\text{deviceChange} \text{ EXCEPT} \\ & \quad \quad ![d] = [\text{deviceChange}[d] \text{ EXCEPT } ![c].\text{state} = \text{Complete}]] \\ & \vee \wedge \text{deviceState}[d] = \text{Disconnected} \\ & \quad \wedge \text{deviceChange}' = [\text{deviceChange} \text{ EXCEPT} \\ & \quad \quad ![d] = [\text{deviceChange}[d] \text{ EXCEPT } ![c].\text{state} = \text{Failed}]] \\ & \wedge \text{UNCHANGED } \langle \text{nodeVars}, \text{networkChange}, \text{deviceVars}, \text{constraintVars} \rangle \end{aligned}$$

This section models device states. Devices begin in the *Disconnected* state and can only be configured while in the *Connected* state.

Set device *d* state to *Connected*

$$\begin{aligned} \text{ConnectDevice}(d) & \triangleq \\ & \wedge \text{deviceState}' = [\text{deviceState} \text{ EXCEPT } ![d] = \text{Connected}] \\ & \wedge \text{connectionCount}' = \text{connectionCount} + 1 \\ & \wedge \text{UNCHANGED } \langle \text{nodeVars}, \text{configVars}, \text{electionCount}, \text{configCount} \rangle \end{aligned}$$

Set device *d* state to *Disconnected*

$$\begin{aligned} \text{DisconnectDevice}(d) & \triangleq \\ & \wedge \text{deviceState}' = [\text{deviceState} \text{ EXCEPT } ![d] = \text{Disconnected}] \\ & \wedge \text{connectionCount}' = \text{connectionCount} + 1 \\ & \wedge \text{UNCHANGED } \langle \text{nodeVars}, \text{configVars}, \text{electionCount}, \text{configCount} \rangle \end{aligned}$$

Init and next state predicates

$$\begin{aligned} \text{Init} & \triangleq \\ & \wedge \text{leader} = [n \in \text{Node} \mapsto \text{FALSE}] \\ & \wedge \text{master} = [n \in \text{Node} \mapsto [d \in \text{Device} \mapsto \text{FALSE}]] \\ & \wedge \text{networkChange} = \langle \rangle \\ & \wedge \text{deviceChange} = [d \in \text{Device} \mapsto [x \in \{ \} \mapsto [\text{phase} \mapsto \text{Change}, \text{state} \mapsto \text{Pending}]]] \\ & \wedge \text{deviceState} = [d \in \text{Device} \mapsto \text{Disconnected}] \\ & \wedge \text{electionCount} = 0 \\ & \wedge \text{configCount} = 0 \\ & \wedge \text{connectionCount} = 0 \end{aligned}$$

$$\text{Next} \triangleq$$

$$\begin{aligned}
& \forall \exists d \in \text{SUBSET } Device : \\
& \quad SubmitChange([x \in d \mapsto 1]) \\
& \forall \exists c \in \text{DOMAIN } networkChange : \\
& \quad RollbackChange(c) \\
& \forall \exists n \in Node : \\
& \quad \exists l \in Node : \\
& \quad \quad SetNodeLeader(n, l) \\
& \forall \exists n \in Node : \\
& \quad \exists d \in Device : \\
& \quad \quad \exists l \in Node : \\
& \quad \quad \quad SetDeviceMaster(n, d, l) \\
& \forall \exists n \in Node : \\
& \quad \exists c \in \text{DOMAIN } networkChange : \\
& \quad \quad ReconcileNetworkChange(n, c) \\
& \forall \exists n \in Node : \\
& \quad \exists d \in Device : \\
& \quad \quad \exists c \in \text{DOMAIN } deviceChange[d] : \\
& \quad \quad \quad ReconcileNetworkChange(n, c) \\
& \forall \exists n \in Node : \\
& \quad \exists d \in Device : \\
& \quad \quad \exists c \in \text{DOMAIN } deviceChange[d] : \\
& \quad \quad \quad ReconcileDeviceChange(n, d, c) \\
& \forall \exists d \in Device : \\
& \quad ConnectDevice(d) \\
& \forall \exists d \in Device : \\
& \quad DisconnectDevice(d) \\
& Spec \triangleq Init \wedge \Box[Next]_{vars}
\end{aligned}$$

* Modification History
* Last modified Wed Sep 22 18:37:38 PDT 2021 by jordanhalterman
* Created Wed Sep 22 13:22:32 PDT 2021 by jordanhalterman