
MODULE *Config*

INSTANCE *Naturals*

INSTANCE *FiniteSets*

INSTANCE *Sequences*

INSTANCE *TLC*

An empty constant

CONSTANT *Nil*

Transaction type constants

CONSTANTS

Change,
Rollback

Transaction isolation constants

CONSTANTS

ReadCommitted,
Serializable

Phase constants

CONSTANTS

Initialize,
Validate,
Abort,
Commit,
Apply

Phase \triangleq

$\{$ *Initialize*,
Validate,
Abort,
Commit,
Apply $\}$

Status constants

CONSTANTS

InProgress,
Complete,
Failed

State \triangleq

$\{$ *InProgress*,
Complete,
 $\}$

Failed}

State constants

CONSTANTS

Pending,
Validated,
Committed,
Applied,
Aborted

$Status \triangleq$
 $\{Pending,$
 $Validated,$
 $Committed,$
 $Applied,$
 $Aborted\}$

CONSTANTS

Valid,
Invalid

CONSTANTS

Success,
Failure

The set of all nodes

CONSTANT *Node*

Target is the set of all targets and their possible paths and values.

Example:

$Target \triangleq$
 $[target1 \mapsto$
 $\quad [persistent \mapsto FALSE, values \mapsto [$
 $\quad \quad path1 \mapsto \{“value1”, “value2”\},$
 $\quad \quad path2 \mapsto \{“value2”, “value3”\}]],$
 $target2 \mapsto$
 $\quad [persistent \mapsto TRUE, values \mapsto [$
 $\quad \quad path2 \mapsto \{“value3”, “value4”\},$
 $\quad \quad path3 \mapsto \{“value4”, “value5”\}]]]$

CONSTANT *Target*

$Empty \triangleq [p \in \{\} \mapsto [value \mapsto Nil, delete \mapsto FALSE]]$

Configuration update/rollback requests are tracked and processed through two data types. Transactions represent the lifecycle of a single configuration change request and are stored in an append-only log. Configurations represent the desired configuration of a *gNMI* target based on the aggregate of relevant changes in the Transaction log.

```

TYPE Type ::= type ∈
  {Change,
   Rollback}

TYPE Phase ::= phase ∈
  {Initialize,
   Validate,
   Abort,
   Commit,
   Apply}

TYPE State ::= state ∈
  {InProgress,
   Complete,
   Failed}

TYPE Status ::= status ∈
  {Pending,
   Validated,
   Committed,
   Applied,
   Aborted}

TYPE Isolation ::= isolation ∈
  {ReadCommitted,
   Serializable}

TYPE Transaction  $\triangleq$ 
  [type      ::= type ∈ Type,
   isolation ::= isolation ∈ Isolation
   change ::=
     [target ∈ SUBSET (DOMAIN Target) ↦
      [path ∈ SUBSET (DOMAIN Target[target].values) ↦
       [value ::= value ∈ STRING,
        delete ::= delete ∈ BOOLEAN ]]],
   rollback ::= index ∈ Nat,
   targets ::= targets ∈ SUBSET (DOMAIN Target)
   phase     ::= phase ∈ Phase,
   state     ::= state ∈ State,
   status    ::= status ∈ Status]

TYPE Proposal  $\triangleq$ 
  [type      ::= type ∈ Type,
   change     ::=
     [index ::= index ∈ Nat,
      values ::=
        [path ∈ SUBSET (DOMAIN Target[target].values) ↦
         [value ::= value ∈ STRING,
          delete ::= delete ∈ BOOLEAN ]]],
   rollback ::=
     [index ::= index ∈ Nat,
      values ::=
        [path ∈ SUBSET (DOMAIN Target[target].values) ↦

```

```

    [value ::= value ∈ STRING,
     delete ::= delete ∈ BOOLEAN ]],
    dependency ::= [index ∈ Nat],
    phase      ::= phase ∈ Phase,
    state      ::= state ∈ State]
TYPE Configuration  $\triangleq$ 
[config ::=
  [index ::= index ∈ Nat,
   term  ::= term ∈ Nat,
   values ::=
    [path ∈ SUBSET (DOMAIN Target[target])  $\mapsto$ 
     [value ::= value ∈ STRING,
      index ::= index ∈ Nat,
      deleted ::= delete ∈ BOOLEAN ]],
   proposal ::= [index ::= index ∈ Nat],
   commit ::= [index ::= index ∈ Nat],
   target ::=
    [index ::= index ∈ Nat,
     term  ::= term ∈ Nat,
     values ::=
      [path ∈ SUBSET (DOMAIN Target[target])  $\mapsto$ 
       [value ::= value ∈ STRING,
        index ::= index ∈ Nat,
        deleted ::= delete ∈ BOOLEAN ]],
     state ::= state ∈ State]

```

A transaction log. Transactions may either request a set of changes to a set of targets or rollback a prior change.

VARIABLE *transaction*

A record of per-target proposals

VARIABLE *proposal*

A record of per-target configurations

VARIABLE *configuration*

A record of target states

VARIABLE *target*

A record of target masterhips

VARIABLE *mastership*

$vars \triangleq \langle transaction, proposal, configuration, mastership, target \rangle$

This section models *mastership* for the configuration service.

Mastership is used primarily to track the lifecycle of individual configuration targets and react to state changes on the southbound. Each target is assigned a master from the *Node* set, and masters can be unset when the target disconnects.

Set node n as the master for target t
 $SetMaster(n, t) \triangleq$
 $\wedge mastership[t].master \neq n$
 $\wedge mastership' = [mastership \text{ EXCEPT } ![t].term = mastership[t].term + 1,$
 $\phantom{\wedge mastership' = [mastership \text{ EXCEPT } } ![t].master = n]$
 $\wedge \text{UNCHANGED } \langle transaction, proposal, configuration, target \rangle$
 $UnsetMaster(t) \triangleq$
 $\wedge mastership[t].master \neq Nil$
 $\wedge mastership' = [mastership \text{ EXCEPT } ![t].master = Nil]$
 $\wedge \text{UNCHANGED } \langle transaction, proposal, configuration, target \rangle$

This section models configuration changes and rollbacks. Changes are appended to the transaction log and processed asynchronously.

$Value(s, t, p) \triangleq$
 $\text{LET } value \triangleq \text{CHOOSE } v \in s : v.target = t \wedge v.path = p$
 IN
 $[value \mapsto value.value,$
 $ delete \mapsto value.delete]$
 $Paths(s, t) \triangleq$
 $[p \in \{v.path : v \in \{v \in s : v.target = t\}\} \mapsto Value(s, t, p)]$
 $Changes(s) \triangleq$
 $[t \in \{v.target : v \in s\} \mapsto Paths(s, t)]$
 $ValidValues(t, p) \triangleq$
 $\text{UNION } \{[value \mapsto v, delete \mapsto \text{FALSE}] : v \in Target[t].values[p], [value \mapsto Nil, delete \mapsto \text{TRUE}]\}$
 $ValidPaths(t) \triangleq$
 $\text{UNION } \{[v @@@ [path \mapsto p] : v \in ValidValues(t, p)] : p \in \text{DOMAIN } Target[t].values\}$
 $ValidTargets \triangleq$
 $\text{UNION } \{[p @@@ [target \mapsto t] : p \in ValidPaths(t)] : t \in \text{DOMAIN } Target\}$

The set of all valid sets of changes to all targets and their paths.

The set of possible changes is computed from the *Target* model value.

$ValidChanges \triangleq$
 $\text{LET } changeSets \triangleq \{s \in \text{SUBSET } ValidTargets :$
 $\phantom{\text{LET } changeSets \triangleq } \forall t \in \text{DOMAIN } Target :$
 $\phantom{\text{LET } changeSets \triangleq } \forall p \in \text{DOMAIN } Target[t].values :$
 $\phantom{\text{LET } changeSets \triangleq } Cardinality(\{v \in s : v.target = t \wedge v.path = p\}) \leq 1\}$
 IN
 $\{Changes(s) : s \in changeSets\}$

The next available index in the transaction log.

This is computed as the max of the existing indexes in the log to allow for changes to the log (*e.g.* log compaction) to be modeled.

$NextIndex \triangleq$

```

IF DOMAIN transaction = {} THEN
  1
ELSE
  LET i  $\triangleq$  CHOOSE i ∈ DOMAIN transaction :
    ∀ j ∈ DOMAIN transaction : i ≥ j
  IN i + 1

```

Add a set of changes 'c' to the transaction log

$RequestChange(c) \triangleq$

```

∧ ∃ isolation ∈ {ReadCommitted, Serializable} :
  ∧ transaction' = transaction @@ (NextIndex :> [
    type      ↦ Change,
    isolation ↦ isolation,
    change    ↦ c,
    targets   ↦ {},
    phase     ↦ Initialize,
    state     ↦ InProgress,
    status    ↦ Pending])
  ∧ UNCHANGED ⟨proposal, configuration, mastership, target⟩

```

Add a rollback of transaction 't' to the transaction log

$RequestRollback(i) \triangleq$

```

∧ ∃ isolation ∈ {ReadCommitted, Serializable} :
  ∧ transaction' = transaction @@ (NextIndex :> [
    type      ↦ Rollback,
    isolation ↦ isolation,
    rollback  ↦ i,
    targets   ↦ {},
    phase     ↦ Initialize,
    state     ↦ InProgress,
    status    ↦ Pending])
  ∧ UNCHANGED ⟨proposal, configuration, mastership, target⟩

```

This section models the Transaction log reconciler.

Transactions come in two flavors : – *Change* transactions contain a set of changes to be applied to a set of *targets* – *Rollback* transactions reference a prior change transaction to be reverted to the previous state

Transactions proceed through a series of phases:

- * *Initialize* – create and link Proposals
- * *Validate* – validate changes and rollbacks
- * *Commit* – *commit* changes to Configurations
- * *Apply* – *commit* changes to Targets

Reconcile a transaction
 $ReconcileTransaction(n, i) \triangleq$

Initialize is the only transaction phase that's globally serialized.
While in the *Initializing* phase, the reconciler checks whether the prior transaction has been *Initialized* before creating Proposals in the *Initialize* phase. Once all of the transaction's proposals have been *Initialized*, the transaction will be marked *Initialized*. If any proposal is *Failed*, the transaction will be marked *Failed* as well.

$\wedge \vee \wedge transaction[i].phase = Initialize$
 $\wedge \vee \wedge transaction[i].state = InProgress$

All prior transaction must be initialized before proceeding to initialize this transaction.

$\wedge \neg \exists j \in \text{DOMAIN } transaction :$
 $\wedge j < i$
 $\wedge transaction[j].phase = Initialize$
 $\wedge transaction[j].state = InProgress$

If the transaction's targets are not yet set, create proposals and add targets to the transaction state.

$\wedge \vee \wedge transaction[i].targets = \{\}$

If the transaction is a change, the targets are taken from the change values.

$\wedge \vee \wedge transaction[i].type = Change$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].targets = \text{DOMAIN } transaction[i].change]$
 $\wedge proposal' = [t \in \text{DOMAIN } proposal \mapsto$
 IF $t \in \text{DOMAIN } transaction[i].change$ THEN
 $proposal[t] @@ (i :> [type \mapsto Change,$
 $change \mapsto$
 $[index \mapsto i,$
 $values \mapsto transaction[i].change[t]],$
 $rollback \mapsto$
 $[index \mapsto 0,$
 $values \mapsto Empty],$
 $dependency \mapsto [index \mapsto 0],$
 $phase \mapsto Initialize,$
 $state \mapsto InProgress])$
 ELSE
 $proposal[t]$

If the transaction is a rollback, the targets affected are the targets of the change transaction being rolled back.

$\vee \wedge transaction[i].type = Rollback$

If the rollback index is a valid *Change* transaction, initialize proposals for all of the *Change* targets.

$\wedge \vee \wedge transaction[i].rollback \in \text{DOMAIN } transaction$
 $\wedge transaction[transaction[i].rollback].type = Change$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].targets =$

DOMAIN $transaction[transaction[i].rollback].change]$

$\wedge proposal' = [t \in \text{DOMAIN } proposal \mapsto$

IF $t \in \text{DOMAIN } transaction[transaction[i].rollback].change$ THEN

$proposal[t] @@ (i :> [type \mapsto Rollback,$

$change \mapsto$

$[index \mapsto 0,$

$values \mapsto Empty],$

$rollback \mapsto$

$[index \mapsto transaction[i].rollback,$

$values \mapsto Empty],$

$dependency \mapsto [index \mapsto 0],$

$phase \mapsto Initialize,$

$state \mapsto InProgress])$

ELSE

$proposal[t]$

If the rollback index is not a valid *Change* transaction

fail the *Rollback* transaction.

$\vee \wedge \vee \wedge transaction[i].rollback \in \text{DOMAIN } transaction$

$\wedge transaction[transaction[i].rollback].type = Rollback$

$\vee transaction[i].rollback \notin \text{DOMAIN } transaction$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Failed]$

$\wedge \text{UNCHANGED } \langle proposal \rangle$

If the transaction's proposals have been initialized, check proposals

for completion or failures.

$\vee \wedge transaction[i].targets \neq \{\}$

If all proposals have been *Complete*, mark the transaction *Complete*.

$\wedge \vee \wedge \forall t \in transaction[i].targets :$

$\wedge proposal[t][i].phase = Initialize$

$\wedge proposal[t][i].state = Complete$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Complete]$

$\wedge \text{UNCHANGED } \langle proposal \rangle$

If any proposal has been *Failed*, mark the transaction *Failed*.

$\vee \wedge \exists t \in transaction[i].targets :$

$\wedge proposal[t][i].phase = Initialize$

$\wedge proposal[t][i].state = Failed$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Failed]$

$\wedge \text{UNCHANGED } \langle proposal \rangle$

Once the transaction has been *Initialized*, proceed to the *Validate* phase.

If any of the transaction's proposals depend on a *Serializable* transaction,

verify the dependency has been *Validated* to preserve serializability before

moving the transaction to the *Validate* phase.

$\vee \wedge transaction[i].state = Complete$

$\wedge \forall t \in transaction[i].targets :$

$\wedge proposal[t][i].dependency.index \in \text{DOMAIN } transaction$

$\wedge transaction[proposal[t][i].dependency.index].isolation = Serializable$

$\Rightarrow transaction[proposal[t][i].dependency.index].status \in \{Validated, Committed, Applied, A\}$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Validate,$
 $![i].state = InProgress]$
 $\wedge \text{UNCHANGED } \langle proposal \rangle$
 If the transaction failed initialization, proceed to the *Abort* phase
 to ensure indexes are still updated for the target configurations.
 $\vee \wedge transaction[i].state = Failed$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Abort,$
 $![i].state = InProgress]$
 $\wedge \text{UNCHANGED } \langle proposal \rangle$
 $\vee \wedge transaction[i].phase = Validate$
 $\wedge \vee \wedge transaction[i].state = InProgress$
 Move the transaction's proposals to the *Validating* state
 $\wedge \vee \wedge \exists t \in transaction[i].targets :$
 $\wedge proposal[t][i].phase \neq Validate$
 $\wedge proposal' = [proposal \text{ EXCEPT } ![t] =$
 $[proposal[t] \text{ EXCEPT } ![i].phase = Validate,$
 $![i].state = InProgress]]$
 $\wedge \text{UNCHANGED } \langle transaction \rangle$
 If all proposals have been *Complete*, mark the transaction *Complete*.
 $\vee \wedge \forall t \in transaction[i].targets :$
 $\wedge proposal[t][i].phase = Validate$
 $\wedge proposal[t][i].state = Complete$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Complete,$
 $![i].status = Validated]$
 $\wedge \text{UNCHANGED } \langle proposal \rangle$
 If any proposal has been *Failed*, mark the transaction *Failed*.
 $\vee \wedge \exists t \in transaction[i].targets :$
 $\wedge proposal[t][i].phase = Validate$
 $\wedge proposal[t][i].state = Failed$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Failed]$
 $\wedge \text{UNCHANGED } \langle proposal \rangle$
 Once the transaction has been *Validated*, proceed to the *Commit* phase.
 If any of the transaction's proposals depend on a *Serializable* transaction,
 verify the dependency has been *Committed* to preserve serializability before
 moving the transaction to the *Commit* phase.
 $\vee \wedge transaction[i].state = Complete$
 $\wedge \forall t \in transaction[i].targets :$
 $\wedge proposal[t][i].dependency.index \in \text{DOMAIN } transaction$
 $\wedge transaction[proposal[t][i].dependency.index].isolation = Serializable$
 $\Rightarrow transaction[proposal[t][i].dependency.index].status \in \{Committed, Applied, Aborted\}$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Commit,$
 $![i].state = InProgress]$
 $\wedge \text{UNCHANGED } \langle proposal \rangle$
 If the transaction failed validation, proceed to the *Abort* phase

to ensure indexes are still updated for the target configurations.

$$\begin{aligned} & \vee \wedge \text{transaction}[i].\text{state} = \text{Failed} \\ & \wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{phase} = \text{Abort}, \\ & \hspace{15em} ![i].\text{state} = \text{InProgress}] \\ & \wedge \text{UNCHANGED } \langle \text{proposal} \rangle \\ & \vee \wedge \text{transaction}[i].\text{phase} = \text{Commit} \\ & \wedge \vee \wedge \text{transaction}[i].\text{state} = \text{InProgress} \\ & \quad \text{Move the transaction's proposals to the } \textit{Committing} \text{ state} \\ & \wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : \\ & \quad \wedge \text{proposal}[t][i].\text{phase} \neq \text{Commit} \\ & \quad \wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = \\ & \hspace{15em} [\text{proposal}[t] \text{ EXCEPT } ![i].\text{phase} = \text{Commit}, \\ & \hspace{15em} ![i].\text{state} = \text{InProgress}]] \\ & \wedge \text{UNCHANGED } \langle \text{transaction} \rangle \\ & \quad \text{If all proposals have been } \textit{Complete}, \text{ mark the transaction } \textit{Complete}. \\ & \vee \wedge \forall t \in \text{transaction}[i].\text{targets} : \\ & \quad \wedge \text{proposal}[t][i].\text{phase} = \text{Commit} \\ & \quad \wedge \text{proposal}[t][i].\text{state} = \text{Complete} \\ & \quad \wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{state} = \text{Complete}, \\ & \hspace{15em} ![i].\text{status} = \text{Committed}] \\ & \wedge \text{UNCHANGED } \langle \text{proposal} \rangle \\ & \quad \text{Once the transaction has been } \textit{Committed}, \text{ proceed to the } \textit{Apply} \text{ phase.} \\ & \quad \text{If any of the transaction's proposals depend on a } \textit{Serializable} \text{ transaction,} \\ & \quad \text{verify the dependency has been } \textit{Applied} \text{ to preserve serializability before} \\ & \quad \text{moving the transaction to the } \textit{Apply} \text{ phase.} \\ & \vee \wedge \text{transaction}[i].\text{state} = \text{Complete} \\ & \quad \wedge \forall t \in \text{transaction}[i].\text{targets} : \\ & \quad \quad \wedge \text{proposal}[t][i].\text{dependency.index} \in \text{DOMAIN } \text{transaction} \\ & \quad \quad \wedge \text{transaction}[\text{proposal}[t][i].\text{dependency.index}].\text{isolation} = \textit{Serializable} \\ & \quad \quad \Rightarrow \text{transaction}[\text{proposal}[t][i].\text{dependency.index}].\text{status} \in \{\textit{Applied}, \textit{Aborted}\} \\ & \quad \wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{phase} = \textit{Apply}, \\ & \hspace{15em} ![i].\text{state} = \textit{InProgress}] \\ & \wedge \text{UNCHANGED } \langle \text{proposal} \rangle \\ & \vee \wedge \text{transaction}[i].\text{phase} = \textit{Apply} \\ & \quad \wedge \text{transaction}[i].\text{state} = \textit{InProgress} \\ & \quad \text{Move the transaction's proposals to the } \textit{Applying} \text{ state} \\ & \wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : \\ & \quad \wedge \text{proposal}[t][i].\text{phase} \neq \textit{Apply} \\ & \quad \wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = \\ & \hspace{15em} [\text{proposal}[t] \text{ EXCEPT } ![i].\text{phase} = \textit{Apply}, \\ & \hspace{15em} ![i].\text{state} = \textit{InProgress}]] \\ & \wedge \text{UNCHANGED } \langle \text{transaction} \rangle \\ & \quad \text{If all proposals have been } \textit{Complete}, \text{ mark the transaction } \textit{Complete}. \\ & \vee \wedge \forall t \in \text{transaction}[i].\text{targets} : \\ & \quad \wedge \text{proposal}[t][i].\text{phase} = \textit{Apply} \end{aligned}$$

$\wedge \text{proposal}[t][i].\text{state} = \text{Complete}$
 $\wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{state} = \text{Complete},$
 $\phantom{\wedge \text{transaction}' = [} ![i].\text{status} = \text{Applied}]$
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$
 If any proposal has been *Failed*, mark the transaction *Failed*.
 $\vee \wedge \exists t \in \text{transaction}[i].\text{targets} :$
 $\phantom{\vee \wedge \exists t \in \text{transaction}[i].\text{targets} : } \wedge \text{proposal}[t][i].\text{phase} = \text{Apply}$
 $\phantom{\vee \wedge \exists t \in \text{transaction}[i].\text{targets} : } \wedge \text{proposal}[t][i].\text{state} = \text{Failed}$
 $\wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{state} = \text{Failed}]$
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$
 The *Aborting* state is used to clean up transactions that have failed during
 the *Initializing* or *Validating* phases.
 $\vee \wedge \text{transaction}[i].\text{phase} = \text{Abort}$
 $\wedge \text{transaction}[i].\text{state} = \text{InProgress}$
 Move the transaction's proposals to the *Aborting* state
 $\wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} :$
 $\phantom{\wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : } \wedge \text{proposal}[t][i].\text{phase} \neq \text{Abort}$
 $\phantom{\wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : } \wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] =$
 $\phantom{\wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : } \phantom{\wedge \text{proposal}' = [} [\text{proposal}[t] \text{ EXCEPT } ![i].\text{phase} = \text{Abort},$
 $\phantom{\wedge \vee \wedge \exists t \in \text{transaction}[i].\text{targets} : } \phantom{\wedge \text{proposal}' = [} ![i].\text{state} = \text{InProgress}]]$
 $\wedge \text{UNCHANGED } \langle \text{transaction} \rangle$
 If all proposals have been *Complete*, mark the transaction *Complete*.
 $\vee \wedge \forall t \in \text{transaction}[i].\text{targets} :$
 $\phantom{\vee \wedge \forall t \in \text{transaction}[i].\text{targets} : } \wedge \text{proposal}[t][i].\text{phase} = \text{Abort}$
 $\phantom{\vee \wedge \forall t \in \text{transaction}[i].\text{targets} : } \wedge \text{proposal}[t][i].\text{state} = \text{Complete}$
 $\wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{state} = \text{Complete},$
 $\phantom{\wedge \text{transaction}' = [} ![i].\text{status} = \text{Aborted}]$
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$
 $\wedge \text{UNCHANGED } \langle \text{configuration}, \text{mastership}, \text{target} \rangle$
 Reconcile a proposal
 $\text{ReconcileProposal}(n, t, i) \triangleq$
 $\wedge \vee \wedge \text{proposal}[t][i].\text{phase} = \text{Initialize}$
 $\phantom{\wedge \vee \wedge \text{proposal}[t][i].\text{phase} = \text{Initialize} } \wedge \text{proposal}[t][i].\text{state} = \text{InProgress}$
 $\phantom{\wedge \vee \wedge \text{proposal}[t][i].\text{phase} = \text{Initialize} } \wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] =$
 $\phantom{\wedge \vee \wedge \text{proposal}[t][i].\text{phase} = \text{Initialize} } \phantom{\wedge \text{proposal}' = [} [\text{proposal}[t] \text{ EXCEPT } ![i].\text{state} = \text{Complete},$
 $\phantom{\wedge \vee \wedge \text{proposal}[t][i].\text{phase} = \text{Initialize} } \phantom{\wedge \text{proposal}' = [} ![i].\text{dependency.index} = \text{configuration}[t].\text{proposal.index}]$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{proposal.index} = i]$
 $\wedge \text{UNCHANGED } \langle \text{target} \rangle$
 While in the *Validate* phase, validate the proposed changes.
 If validation is successful, the proposal also records the changes
 required to roll back the proposal and the index to which to roll back.
 $\vee \wedge \text{proposal}[t][i].\text{phase} = \text{Validate}$
 $\phantom{\vee \wedge \text{proposal}[t][i].\text{phase} = \text{Validate} } \wedge \text{proposal}[t][i].\text{state} = \text{InProgress}$
 $\phantom{\vee \wedge \text{proposal}[t][i].\text{phase} = \text{Validate} } \wedge \text{configuration}[t].\text{commit.index} = \text{proposal}[t][i].\text{dependency.index}$

$$\begin{aligned} & \wedge \vee \wedge \text{proposal}[t][i].\text{type} = \text{Change} \\ & \wedge \text{LET } \text{rollbackIndex} \triangleq \text{configuration}[t].\text{config.index} \\ & \quad \text{rollbackValues} \triangleq [p \in \text{DOMAIN } \text{proposal}[t][i].\text{change.values} \mapsto \\ & \quad \quad \text{IF } p \in \text{DOMAIN } \text{configuration}[t].\text{config.values} \text{ THEN} \\ & \quad \quad \quad \text{configuration}[t].\text{config.values}[p] \\ & \quad \quad \text{ELSE} \\ & \quad \quad \quad [\text{value} \mapsto \text{Nil}, \\ & \quad \quad \quad \text{delete} \mapsto \text{TRUE}]] \end{aligned}$$
$$\text{IN} \quad \exists r \in \{Valid, Invalid\} :$$
$$\wedge \textit{proposal}' = [\textit{proposal} \text{ EXCEPT } ![t] = \\ \text{[proposal}[t] \text{ EXCEPT } ![i].\textit{rollback.index} = \textit{rollbackIndex}, \\ [t] \text{ EXCEPT } } ![i].\textit{rollback.values} = \textit{rollbackValues}, \\ [t] \text{ EXCEPT } } ![i].\textit{state} = \textit{Complete}]]$$
$$\wedge proposal' = [proposal \text{ EXCEPT } ![t] = [proposal[t] \text{ EXCEPT } ![i].state = Failed]]$$
$$\wedge proposal' = [proposal \text{ EXCEPT } ![t] =$$

$[proposal[t] \text{ EXCEPT } ![i].state = Failed]$

If the *Rollback* target is not the most recent change to the configuration,
fail validation for the proposal.
 $\vee \wedge configuration[t].config.index \neq proposal[t][i].rollback.index$
 $\wedge proposal' = [proposal \text{ EXCEPT } ![t] = [proposal[t] \text{ EXCEPT } ![i].state = Failed]]$

If a *Rollback* proposal is attempting to roll back another *Rollback*,
fail validation for the proposal.
 $\vee \wedge proposal[t][proposal[t][i].rollback.index].type = Rollback$
 $\wedge proposal' = [proposal \text{ EXCEPT } ![t] =$
 $\quad [proposal[t] \text{ EXCEPT } ![i].state = Failed]]$

$\wedge \text{UNCHANGED } \langle configuration, target \rangle$

While in the *Commit* state, commit the proposed changes to the configuration.
 $\vee \wedge proposal[t][i].phase = Commit$
 $\wedge proposal[t][i].state = InProgress$
Only commit the proposal if the prior proposal has already been committed.
 $\wedge configuration[t].commit.index = proposal[t][i].dependency.index$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![t].config.values = proposal[t][i].change.values,$
 $\quad \quad \quad ![t].config.index = proposal[t][i].change.index,$
 $\quad \quad \quad ![t].commit.index = i]$
 $\wedge proposal' = [proposal \text{ EXCEPT } ![t] = [proposal[t] \text{ EXCEPT } ![i].state = Complete]]$
 $\wedge \text{UNCHANGED } \langle target \rangle$

While in the *Apply* phase, apply the proposed changes to the target.
 $\vee \wedge proposal[t][i].phase = Apply$
 $\wedge proposal[t][i].state = InProgress$
 $\wedge configuration[t].target.index = proposal[t][i].dependency.index$
 $\wedge configuration[t].target.term = mastership[t].term$
 $\wedge mastership[t].master = n$
Model successful and failed target update requests.
 $\wedge \exists r \in \{Success, Failure\} :$
 $\quad \vee \wedge r = Success$
 $\quad \wedge target' = [target \text{ EXCEPT } ![t] = proposal[t][i].change.values @@ target[t]]$
 $\quad \wedge configuration' = [configuration \text{ EXCEPT }$
 $\quad \quad \quad ![t].target.index = i,$
 $\quad \quad \quad ![t].target.values = proposal[t][i].change.values$
 $\quad \quad \quad @@ configuration[t].target.values]$
 $\quad \wedge proposal' = [proposal \text{ EXCEPT } ![t] = [proposal[t] \text{ EXCEPT } ![i].state = Complete]]$
If the proposal could not be applied, update the configuration's applied index
and mark the proposal *Failed*.
 $\vee \wedge r = Failure$
 $\quad \wedge configuration' = [configuration \text{ EXCEPT } ![t].target.index = i]$
 $\quad \wedge proposal' = [proposal \text{ EXCEPT } ![t] = [proposal[t] \text{ EXCEPT } ![i].state = Failed]]$
 $\quad \wedge \text{UNCHANGED } \langle target \rangle$

$\vee \wedge proposal[t][i].phase = Abort$
 $\wedge proposal[t][i].state = InProgress$
The *commit.index* will always be greater than or equal to the *target.index*.

If only the *commit.index* matches the proposal's *dependency.index*, update the *commit.index* to enable commits of later proposals, but do not mark the *Abort* phase *Complete* until the *target.index* has been incremented.
 $\wedge \vee \wedge \text{configuration}[t].\text{commit.index} = \text{proposal}[t][i].\text{dependency.index}$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{commit.index} = i]$
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$
 If the configuration's *target.index* matches the proposal's *dependency.index*, update the *target.index* and mark the proposal *Complete* for the *Abort* phase.
 $\vee \wedge \text{configuration}[t].\text{commit.index} \geq i$
 $\wedge \text{configuration}[t].\text{target.index} = \text{proposal}[t][i].\text{dependency.index}$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{target.index} = i]$
 $\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = [\text{proposal}[t] \text{ EXCEPT } ![i].\text{state} = \text{Complete}]]$
 If both the configuration's *commit.index* and *target.index* match the proposal's *dependency.index*, update the *commit.index* and *target.index* and mark the proposal *Complete* for the *Abort* phase.
 $\vee \wedge \text{configuration}[t].\text{commit.index} = \text{proposal}[t][i].\text{dependency.index}$
 $\wedge \text{configuration}[t].\text{target.index} = \text{proposal}[t][i].\text{dependency.index}$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{commit.index} = i,$
 $\phantom{\wedge \text{configuration}' = [} ![t].\text{target.index} = i]$
 $\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![t] = [\text{proposal}[t] \text{ EXCEPT } ![i].\text{state} = \text{Complete}]]$
 $\wedge \text{UNCHANGED } \langle \text{target} \rangle$
 $\wedge \text{UNCHANGED } \langle \text{transaction}, \text{mastership} \rangle$

This section models the Configuration reconciler.

$\text{ReconcileConfiguration}(n, t) \triangleq$
 $\wedge \vee \wedge \text{Target}[t].\text{persistent}$
 $\wedge \text{configuration}[t].\text{state} \neq \text{Complete}$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{state} = \text{Complete}]$
 $\wedge \text{UNCHANGED } \langle \text{target} \rangle$
 $\vee \wedge \neg \text{Target}[t].\text{persistent}$
 $\wedge \vee \text{mastership}[t].\text{term} > \text{configuration}[t].\text{config.term}$
 $\vee \wedge \text{mastership}[t].\text{term} = \text{configuration}[t].\text{config.term}$
 $\wedge \text{mastership}[t].\text{master} = \text{Nil}$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{config.term} = \text{mastership}[t].\text{term},$
 $\phantom{\wedge \text{configuration}' = [} ![t].\text{state} = \text{InProgress}]$
 $\wedge \text{UNCHANGED } \langle \text{target} \rangle$
 $\vee \wedge \text{configuration}[t].\text{state} = \text{InProgress}$
 $\wedge \text{mastership}[t].\text{term} = \text{configuration}[t].\text{config.term}$
 $\wedge \text{mastership}[t].\text{master} = n$
 $\wedge \text{target}' = [\text{target} \text{ EXCEPT } ![t] = \text{configuration}[t].\text{target.values}]$
 $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } ![t].\text{target.term} = \text{mastership}[t].\text{term},$
 $\phantom{\wedge \text{configuration}' = [} ![t].\text{state} = \text{Complete}]$
 $\wedge \text{UNCHANGED } \langle \text{proposal}, \text{transaction}, \text{mastership} \rangle$

Formal specification, constraints, and theorems.

$Init \triangleq$

$$\begin{aligned}
& \wedge transaction = [i \in \{\} \mapsto \\
& \quad [type \mapsto Change, \\
& \quad \quad phase \mapsto Initialize, \\
& \quad \quad state \mapsto InProgress, \\
& \quad \quad status \mapsto Pending]] \\
& \wedge proposal = [t \in \text{DOMAIN } Target \mapsto \\
& \quad [i \in \{\} \mapsto \\
& \quad \quad [phase \mapsto Initialize, \\
& \quad \quad \quad state \mapsto InProgress]]] \\
& \wedge configuration = [t \in \text{DOMAIN } Target \mapsto \\
& \quad [state \mapsto InProgress, \\
& \quad \quad config \mapsto \\
& \quad \quad \quad [index \mapsto 0, \\
& \quad \quad \quad \quad term \mapsto 0, \\
& \quad \quad \quad \quad values \mapsto \\
& \quad \quad \quad \quad \quad [path \in \{\} \mapsto \\
& \quad \quad \quad \quad \quad \quad [path \mapsto path, \\
& \quad \quad \quad \quad \quad \quad \quad value \mapsto Nil, \\
& \quad \quad \quad \quad \quad \quad \quad index \mapsto 0, \\
& \quad \quad \quad \quad \quad \quad \quad deleted \mapsto FALSE]]], \\
& \quad \quad \quad proposal \mapsto [index \mapsto 0], \\
& \quad \quad \quad commit \mapsto [index \mapsto 0], \\
& \quad \quad \quad target \mapsto \\
& \quad \quad \quad \quad [index \mapsto 0, \\
& \quad \quad \quad \quad \quad term \mapsto 0, \\
& \quad \quad \quad \quad \quad values \mapsto \\
& \quad \quad \quad \quad \quad \quad [path \in \{\} \mapsto \\
& \quad \quad \quad \quad \quad \quad \quad [path \mapsto path, \\
& \quad \quad \quad \quad \quad \quad \quad \quad value \mapsto Nil, \\
& \quad \quad \quad \quad \quad \quad \quad \quad index \mapsto 0, \\
& \quad \quad \quad \quad \quad \quad \quad \quad deleted \mapsto FALSE]]]]]] \\
& \wedge target = [t \in \text{DOMAIN } Target \mapsto \\
& \quad [path \in \{\} \mapsto \\
& \quad \quad [value \mapsto Nil]]] \\
& \wedge mastership = [t \in \text{DOMAIN } Target \mapsto [master \mapsto Nil, term \mapsto 0]]
\end{aligned}$$

$Next \triangleq$

$$\begin{aligned}
& \vee \exists c \in ValidChanges : \\
& \quad RequestChange(c) \\
& \vee \exists t \in \text{DOMAIN } transaction : \\
& \quad RequestRollback(t)
\end{aligned}$$

$$\begin{aligned}
& \forall \exists n \in \text{Node} : \\
& \quad \exists t \in \text{DOMAIN } \text{Target} : \\
& \quad \quad \text{SetMaster}(n, t) \\
& \quad \forall \exists t \in \text{DOMAIN } \text{Target} : \\
& \quad \quad \text{UnsetMaster}(t) \\
& \forall \exists n \in \text{Node} : \\
& \quad \exists t \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \text{ReconcileTransaction}(n, t) \\
& \forall \exists n \in \text{Node} : \\
& \quad \exists t \in \text{DOMAIN } \text{proposal} : \\
& \quad \quad \exists i \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \quad \text{ReconcileProposal}(n, t, i) \\
& \forall \exists n \in \text{Node} : \\
& \quad \exists c \in \text{DOMAIN } \text{configuration} : \\
& \quad \quad \text{ReconcileConfiguration}(n, c) \\
\text{Spec} & \triangleq \text{Init} \wedge \square[\text{Next}]_{\text{vars}} \wedge \text{WF}_{\text{vars}}(\text{Next}) \\
\text{Order} & \triangleq \\
& \forall t \in \text{DOMAIN } \text{proposal} : \\
& \quad \forall i \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \wedge \wedge \text{proposal}[t][i].\text{phase} = \text{Commit} \\
& \quad \quad \wedge \text{proposal}[t][i].\text{state} = \text{InProgress} \\
& \quad \quad \Rightarrow \neg \exists j \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \quad \wedge j > i \\
& \quad \quad \quad \wedge \text{proposal}[t][j].\text{phase} = \text{Commit} \\
& \quad \quad \quad \wedge \text{proposal}[t][j].\text{state} = \text{Complete} \\
& \quad \wedge \wedge \text{proposal}[t][i].\text{phase} = \text{Apply} \\
& \quad \quad \wedge \text{proposal}[t][i].\text{state} = \text{InProgress} \\
& \quad \quad \Rightarrow \neg \exists j \in \text{DOMAIN } \text{proposal}[t] : \\
& \quad \quad \quad \wedge j > i \\
& \quad \quad \quad \wedge \text{proposal}[t][j].\text{phase} = \text{Apply} \\
& \quad \quad \quad \wedge \text{proposal}[t][j].\text{state} = \text{Complete} \\
\text{Consistency} & \triangleq \\
& \forall t \in \text{DOMAIN } \text{target} : \\
& \quad \text{LET} \\
& \quad \quad \text{Compute the transaction indexes that have been applied to the target} \\
& \quad \text{targetIndexes} \triangleq \{i \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \quad \wedge i \in \text{DOMAIN } \text{proposal}[t] \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{phase} = \text{Apply} \\
& \quad \quad \quad \wedge \text{proposal}[t][i].\text{state} = \text{Complete} \\
& \quad \quad \quad \wedge t \in \text{transaction}[i].\text{targets} \\
& \quad \quad \quad \wedge \neg \exists j \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \quad \quad \wedge j > i \\
& \quad \quad \quad \quad \wedge \text{transaction}[j].\text{type} = \text{Rollback}
\end{aligned}$$

$$\begin{aligned}
& \wedge \text{transaction}[j].\text{rollback} = i \\
& \wedge \text{transaction}[j].\text{phase} = \text{Apply} \\
& \wedge \text{transaction}[j].\text{state} = \text{Complete} \} \\
& \text{Compute the set of paths in the target that have been updated by transactions} \\
\text{appliedPaths} & \triangleq \text{UNION } \{ \text{DOMAIN } \text{proposal}[t][i].\text{change.values} : i \in \text{targetIndexes} \} \\
& \text{Compute the highest index applied to the target for each path} \\
\text{pathIndexes} & \triangleq [p \in \text{appliedPaths} \mapsto \text{CHOOSE } i \in \text{targetIndexes} : \\
& \quad \forall j \in \text{targetIndexes} : \\
& \quad \quad \wedge i \geq j \\
& \quad \quad \wedge p \in \text{DOMAIN } \text{proposal}[t][i].\text{change.values}] \\
& \text{Compute the expected target configuration based on the last indexes applied} \\
& \text{to the target for each path.} \\
\text{expectedConfig} & \triangleq [p \in \text{DOMAIN } \text{pathIndexes} \mapsto \text{proposal}[t][\text{pathIndexes}[p]].\text{change.values}[p]] \\
\text{IN} & \\
& \text{target}[t] = \text{expectedConfig} \\
\text{Isolation} & \triangleq \\
& \forall i \in \text{DOMAIN } \text{transaction} : \\
& \quad \wedge \wedge \text{transaction}[i].\text{phase} = \text{Commit} \\
& \quad \wedge \text{transaction}[i].\text{state} = \text{InProgress} \\
& \quad \wedge \text{transaction}[i].\text{isolation} = \text{Serializable} \\
& \quad \Rightarrow \neg \exists j \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \wedge j > i \\
& \quad \quad \wedge \text{transaction}[j].\text{targets} \cap \text{transaction}[i].\text{targets} \neq \{ \} \\
& \quad \quad \wedge \text{transaction}[j].\text{phase} = \text{Commit} \\
& \quad \wedge \wedge \text{transaction}[i].\text{phase} = \text{Apply} \\
& \quad \wedge \text{transaction}[i].\text{state} = \text{InProgress} \\
& \quad \wedge \text{transaction}[i].\text{isolation} = \text{Serializable} \\
& \quad \Rightarrow \neg \exists j \in \text{DOMAIN } \text{transaction} : \\
& \quad \quad \wedge j > i \\
& \quad \quad \wedge \text{transaction}[j].\text{targets} \cap \text{transaction}[i].\text{targets} \neq \{ \} \\
& \quad \quad \wedge \text{transaction}[j].\text{phase} = \text{Apply} \\
\text{Safety} & \triangleq \Box (\text{Order} \wedge \text{Consistency} \wedge \text{Isolation}) \\
\text{THEOREM } \text{Spec} & \Rightarrow \text{Safety} \\
\text{Termination} & \triangleq \\
& \forall i \in \text{DOMAIN } \text{transaction} : \\
& \quad \wedge \text{transaction}[i].\text{phase} \in \{ \text{Apply}, \text{Abort} \} \\
& \quad \wedge \text{transaction}[i].\text{state} = \text{Complete} \\
\text{Liveness} & \triangleq \Box \Diamond \text{Termination} \\
\text{THEOREM } \text{Spec} & \Rightarrow \text{Liveness}
\end{aligned}$$

Type assumptions.

ASSUME $Nil \in \text{STRING}$

ASSUME $\forall phase \in Phase : phase \in \text{STRING}$

ASSUME $\forall state \in State : state \in \text{STRING}$

ASSUME $\forall status \in Status : status \in \text{STRING}$

ASSUME $\wedge IsFiniteSet(Node)$

$\wedge \forall n \in Node :$

$\wedge n \notin \text{DOMAIN } Target$

$\wedge n \in \text{STRING}$

ASSUME $\wedge \forall t \in \text{DOMAIN } Target :$

$\wedge t \notin Node$

$\wedge t \in \text{STRING}$

$\wedge Target[t].persistent \in \text{BOOLEAN}$

$\wedge \forall p \in \text{DOMAIN } Target[t].values :$

$IsFiniteSet(Target[t].values[p])$

\ * Modification History

\ * Last modified *Thu Feb 10 14:42:19 PST 2022* by *jordanhalterman*

\ * Created *Wed Sep 22 13:22:32 PDT 2021* by *jordanhalterman*