
MODULE *Config*

INSTANCE *Naturals*

INSTANCE *FiniteSets*

INSTANCE *Sequences*

INSTANCE *TLC*

An empty constant

CONSTANT *Nil*

Transaction status constants

CONSTANTS

TransactionPending,
TransactionValidating,
TransactionApplying,
TransactionComplete,
TransactionFailed

Configuration status constants

CONSTANTS

ConfigurationPending,
ConfigurationInitializing,
ConfigurationUpdating,
ConfigurationComplete,
ConfigurationFailed

The set of all nodes

CONSTANT *Node*

Target is the possible targets, paths, and values

Example: $Target \triangleq$ [
 $target1 \mapsto$ [
 $path1 \mapsto \{“value1”, “value2”\}$,
 $path2 \mapsto \{“value2”, “value3”\}$],
 $target2 \mapsto$ [
 $path2 \mapsto \{“value3”, “value4”\}$,
 $path3 \mapsto \{“value4”, “value5”\}$]]

CONSTANT *Target*

ASSUME *Nil* ∈ STRING

ASSUME *TransactionPending* ∈ STRING

ASSUME *TransactionValidating* ∈ STRING

ASSUME *TransactionApplying* ∈ STRING

ASSUME $TransactionComplete \in \text{STRING}$
 ASSUME $TransactionFailed \in \text{STRING}$

 ASSUME $ConfigurationPending \in \text{STRING}$
 ASSUME $ConfigurationInitializing \in \text{STRING}$
 ASSUME $ConfigurationUpdating \in \text{STRING}$
 ASSUME $ConfigurationComplete \in \text{STRING}$
 ASSUME $ConfigurationFailed \in \text{STRING}$

 ASSUME $\wedge IsFiniteSet(Node)$
 $\wedge \forall n \in Node :$
 $\wedge n \notin \text{DOMAIN } Target$
 $\wedge n \in \text{STRING}$

 ASSUME $\wedge \forall t \in \text{DOMAIN } Target :$
 $\wedge t \notin Node$
 $\wedge t \in \text{STRING}$
 $\wedge \forall p \in \text{DOMAIN } Target[t] :$
 $IsFiniteSet(Target[t][p])$

TYPE $TransactionStatus ::= status \in$
 $\{TransactionPending,$
 $TransactionValidating,$
 $TransactionApplying,$
 $TransactionComplete,$
 $TransactionFailed\}$

 TYPE $Transaction \triangleq [$
 $id \quad ::= id \in \text{STRING},$
 $index ::= index \in Nat,$
 $revision ::= revision \in Nat,$
 $atomic ::= atomic \in \text{BOOLEAN} ,$
 $sync \quad ::= sync \in \text{BOOLEAN} ,$
 $changes ::= [target \in \text{SUBSET } (\text{DOMAIN } Target) \mapsto [$
 $path \in \text{SUBSET } (\text{DOMAIN } Target[target]) \mapsto [$
 $value ::= value \in \text{STRING},$
 $delete ::= delete \in \text{BOOLEAN }]],$
 $status ::= status \in TransactionStatus]$

 TYPE $ConfigurationStatus ::= status \in$
 $\{ConfigurationPending,$
 $ConfigurationInitializing,$
 $ConfigurationUpdating,$
 $ConfigurationComplete,$
 $ConfigurationFailed\}$

 TYPE $Configuration \triangleq [$
 $id \quad ::= id \in \text{STRING},$
 $revision ::= revision \in Nat,$

```

target ::= target ∈ STRING,
paths ::= [ path ∈ SUBSET (DOMAIN Target[target]) ↦ [
    value ::= value ∈ STRING,
    index ::= index ∈ Nat,
    deleted ::= delete ∈ BOOLEAN ]],
txIndex ::= txIndex ∈ Nat,
syncIndex ::= syncIndex ∈ Nat,
term      ::= term ∈ Nat,
status ::= status ∈ ConfigurationStatus]

```

A sequence of transactions

Each transactions contains a record of 'changes' for a set of targets

VARIABLE *transaction*

A record of target configurations

Each configuration represents the desired state of the target

VARIABLE *configuration*

A record of target states

VARIABLE *targets*

A record of target masters

VARIABLE *masters*

vars \triangleq $\langle transaction, configuration, targets \rangle$

$ChangeMaster(n, t) \triangleq$
 $\wedge masters[t].master \neq n$
 $\wedge masters' = [masters \text{ EXCEPT } ![t].term = masters[t].term + 1,$
 $\phantom{\wedge masters' = [masters \text{ EXCEPT } } ![t].master = n]$
 $\wedge UNCHANGED \langle transaction, configuration \rangle$

This section models the northbound *API* for the configuration service.

This crazy thing returns the set of all possible sets of valid changes

$ValidChanges \triangleq$
 LET $allPaths \triangleq \text{UNION } \{ (DOMAIN \ Target[t]) : t \in DOMAIN \ Target \}$
 $allValues \triangleq \text{UNION } \{ \text{UNION } \{ Target[t][p] : p \in DOMAIN \ Target[t] \} : t \in DOMAIN \ Target \}$
 IN
 $\{ targetPathValues \in SUBSET (Target \times allPaths \times allValues \times BOOLEAN) :$
 $\wedge \forall target \in DOMAIN \ Target :$
 LET $targetIndexes \triangleq \{ i \in 1 \dots Len(targetPathValues) : \wedge targetPathValues[i][1] = target \}$
 IN $\vee Cardinality(targetIndexes) = 0$
 $\vee \wedge Cardinality(targetIndexes) = 1$
 $\wedge \text{LET } targetPathValue \triangleq targetPathValues[\text{CHOOSE } index \in targetIndexes : TRUE]$
 $\phantom{\wedge \text{LET } } targetPath \triangleq targetPathValue[2]$

$$\begin{aligned}
& \text{targetValue} \triangleq \text{targetPathValue}[3] \\
& \text{IN} \\
& \quad \wedge \text{targetPath} \setminus (\text{DOMAIN Target}[\text{target}]) = \{\} \\
& \quad \wedge \text{targetValue} \in \text{Target}[\text{target}][\text{targetPath}] \\
& \text{NextIndex} \triangleq \\
& \quad \text{IF } \text{Len}(\text{transaction}) = 0 \text{ THEN} \\
& \quad \quad 1 \\
& \quad \text{ELSE} \\
& \quad \quad (\text{CHOOSE } i \in \text{DOMAIN transaction} : \\
& \quad \quad \quad \forall j \in \text{DOMAIN transaction} : \\
& \quad \quad \quad \quad \text{transaction}[j].\text{index} < \text{transaction}[i].\text{index}) + 1 \\
& \quad \text{Add a set of changes to the transaction log} \\
& \text{Change} \triangleq \\
& \quad \wedge \exists \text{changes} \in \text{ValidChanges} : \\
& \quad \quad \wedge \text{transaction}' = \text{Append}(\text{transaction}, [\text{index} \mapsto \text{NextIndex}, \\
& \quad \quad \quad \text{atomic} \mapsto \text{FALSE}, \\
& \quad \quad \quad \text{sync} \mapsto \text{FALSE}, \\
& \quad \quad \quad \text{changes} \mapsto \text{changes}, \\
& \quad \quad \quad \text{status} \mapsto \text{TransactionPending}]) \\
& \quad \wedge \text{UNCHANGED } \langle \text{configuration}, \text{targets} \rangle
\end{aligned}$$

This section models the Transaction log reconciler.

$$\begin{aligned}
& \text{Reconcile the transaction log} \\
& \text{ReconcileTransaction}(n, t) \triangleq \\
& \quad \text{If the transaction is } \text{Pending}, \text{ begin validation if the prior transaction} \\
& \quad \text{has already been applied. This simplifies concurrency control in the controller} \\
& \quad \text{and guarantees transactions are applied to the configurations in sequential order.} \\
& \quad \wedge \vee \wedge \text{transaction}[t].\text{status} = \text{TransactionPending} \\
& \quad \quad \wedge \vee \wedge \text{transaction}[t].\text{index} - 1 \in \text{DOMAIN transaction} \\
& \quad \quad \quad \wedge \text{transaction}[\text{transaction}[t].\text{index} - 1].\text{status} \in \{\text{TransactionComplete}, \text{TransactionFailed}\} \\
& \quad \quad \vee \text{transaction}[t].\text{index} - 1 \notin \text{DOMAIN transaction} \\
& \quad \quad \wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![t].\text{status} = \text{TransactionValidating}] \\
& \quad \quad \wedge \text{UNCHANGED } \langle \text{configuration} \rangle \\
& \quad \text{If the transaction is in the } \text{Validating} \text{ state, compute and validate the} \\
& \quad \text{Configuration for each target.} \\
& \quad \vee \wedge \text{transaction}[t].\text{status} = \text{TransactionValidating} \\
& \quad \quad \text{If validation fails any target, mark the transaction } \text{Failed}. \\
& \quad \quad \text{If validation is successful, proceed to } \text{Applying}. \\
& \quad \quad \wedge \exists \text{valid} \in \text{BOOLEAN} : \\
& \quad \quad \quad \vee \wedge \text{valid} \\
& \quad \quad \quad \quad \wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![t].\text{status} = \text{TransactionApplying}] \\
& \quad \quad \quad \vee \wedge \neg \text{valid}
\end{aligned}$$

$\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionFailed]$
 $\wedge \text{UNCHANGED } \langle configuration \rangle$
 If the transaction is in the *Applying* state, update the Configuration for each target and *Complete* the transaction.
 $\vee \wedge transaction[t].status = TransactionApplying$
 $\wedge \vee \wedge transaction[t].atomic$
 TODO: Apply atomic transactions here
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionComplete]$
 $\wedge \text{UNCHANGED } \langle configuration \rangle$
 $\vee \wedge \neg transaction[t].atomic$
 Add the transaction index to each updated path
 $\wedge configuration' = [$
 $\quad r \in \text{DOMAIN } Target \mapsto [$
 $\quad \quad configuration[r] \text{ EXCEPT}$
 $\quad \quad \quad !.paths = [path \in \text{DOMAIN } transaction[t].changes \mapsto$
 $\quad \quad \quad \quad transaction[t].changes[path] @@ [index \mapsto transaction[t].index]]$
 $\quad \quad \quad @@ configuration[t].paths,$
 $\quad \quad \quad !.txIndex = transaction[t].index,$
 $\quad \quad \quad !.status = ConfigurationPending]]$
 $\wedge transaction' = [transaction \text{ EXCEPT } ![t].status = TransactionComplete]$
 $\wedge \text{UNCHANGED } \langle targets \rangle$

This section models the Configuration reconciler.

$ReconcileConfiguration(n, c) \triangleq$

$\wedge \vee \wedge configuration[c].status = ConfigurationPending$
 If the configuration is marked *ConfigurationPending* and mastership has changed (indicated by an increased mastership term), mark the configuration *ConfigurationInitializing* to force full re-synchronization.
 $\wedge \vee \wedge masters[configuration[c].target].term > configuration[c].term$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationInitializing,$
 $\quad \quad \quad ![c].term = masters[configuration[c].target].term]$
 If the configuration is marked *ConfigurationPending* and the values have changed (determined by comparing the transaction index to the last *sync* index), mark the configuration *ConfigurationUpdating* to push the changes to the target.
 $\vee \wedge configuration[c].syncIndex < configuration[c].txIndex$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationUpdating]$
 $\vee \wedge configuration[c].status = ConfigurationInitializing$
 $\wedge masters[configuration[c].target].master = n$
 Merge the configuration paths with the target paths, removing paths that have been marked deleted
 $\wedge \text{LET } deletePaths \triangleq \{p \in \text{DOMAIN } configuration[c].paths : configuration[c].paths[p].deleted\}$
 $\quad \quad \quad configPaths \triangleq \text{DOMAIN } c.paths \setminus deletePaths$

$targetPaths \triangleq \text{DOMAIN } targets[configuration[c].target] \setminus deletePaths$
 IN
 $\wedge targets' = [targets \text{ EXCEPT } ![configuration[c].target] =$
 $\quad [p \in configPaths \mapsto [value \mapsto configuration[c].paths[p]]]$
 $\quad @@ [p \in targetPaths \mapsto targets[configuration[c].target][p]]]$
 Set the configuration's status to *Complete*
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationComplete,$
 $\quad ![c].syncIndex = configuration[c].txIndex]$
 If the configuration is marked *ConfigurationUpdating*, we only need to
 push paths that have changed since the target was initialized or last
 updated by the controller. The set of changes made since the last
 synchronization are identified by comparing the index of each path-value
 to the last synchronization index, *syncIndex*
 $\vee \wedge configuration[c].status = ConfigurationUpdating$
 $\wedge masters[configuration[c].target].master = n$
 Compute the set of updated and deleted paths by comparing
 their indexes to the target's last sync index.
 $\wedge \text{LET } updatePaths \triangleq \{p \in \text{DOMAIN } configuration[c].paths :$
 $\quad configuration[c].paths[p].index > configuration[c].syncIndex\}$
 $\quad deletePaths \triangleq \{p \in updatePaths : configuration[c].paths[p].deleted\}$
 $\quad configPaths \triangleq updatePaths \setminus deletePaths$
 $\quad targetPaths \triangleq \text{DOMAIN } targets[configuration[c].target] \setminus deletePaths$
 IN
 Update the target paths by adding/updating paths that have changed and
 removing paths that have been deleted since the last *sync*.
 $\wedge targets' = [targets \text{ EXCEPT } ![configuration[c].target] =$
 $\quad [p \in configPaths \mapsto configuration[c].paths[p]]]$
 $\quad @@ [p \in targetPaths \mapsto targets[configuration[c].target][p]]]$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationComplete,$
 $\quad ![c].syncIndex = configuration[c].txIndex]$
 If the configuration is not already *ConfigurationPending* and mastership
 has been lost revert it. This can occur when the connection to the
 target has been lost and the mastership is no longer valid.
TODO: We still need to model mastership changes
 $\vee \wedge c.status \neq ConfigurationPending$
 $\wedge masters[configuration[c].target].master = Nil$
 $\wedge configuration' = [configuration \text{ EXCEPT } ![c].status = ConfigurationPending]$
 $\wedge \text{UNCHANGED } \langle transaction \rangle$

Init and next state predicates

$Init \triangleq$
 $\wedge transaction = \langle \rangle$
 $\wedge configuration = [t \in Target \mapsto$

$$\begin{aligned}
& [id \mapsto t, \\
& \quad config \mapsto \\
& \quad \quad [path \in \{\} \mapsto \\
& \quad \quad \quad [path \mapsto path, \\
& \quad \quad \quad \quad value \mapsto Nil, \\
& \quad \quad \quad \quad index \mapsto 0, \\
& \quad \quad \quad \quad deleted \mapsto FALSE]]]] \\
\wedge targets &= [t \in Target \mapsto \\
& \quad [path \in \{\} \mapsto \\
& \quad \quad [value \mapsto Nil]]] \\
\wedge masters &= [t \in Target \mapsto [master \mapsto Nil, term \mapsto 0]] \\
Next &\triangleq \\
& \vee Change \\
& \vee \exists n \in Node : \\
& \quad \vee \exists t \in \text{DOMAIN } Target : \\
& \quad \quad ChangeMaster(n, t) \\
& \quad \vee \exists t \in \text{DOMAIN } transaction : \\
& \quad \quad ReconcileTransaction(n, t) \\
& \quad \vee \exists t \in \text{DOMAIN } configuration : \\
& \quad \quad ReconcileConfiguration(n, t) \\
Spec &\triangleq Init \wedge \Box[Next]_{vars}
\end{aligned}$$

\ * Modification History
\ * Last modified Tue Jan 18 00:24:57 PST 2022 by jordanhalterman
\ * Created Wed Sep 22 13:22:32 PDT 2021 by jordanhalterman