

---

MODULE *Proposal*

---

EXTENDS *Configuration, Mastership*

INSTANCE *Naturals*

INSTANCE *FiniteSets*

LOCAL INSTANCE *TLC*

---

Transaction type constants

CONSTANTS

*ProposalChange,*  
*ProposalRollback*

Phase constants

CONSTANTS

*ProposalCommit,*  
*ProposalApply*

Status constants

CONSTANTS

*ProposalInProgress,*  
*ProposalComplete,*  
*ProposalFailed*

CONSTANT *TraceProposal*

A record of per-target proposals

VARIABLE *proposal*

---

LOCAL *InitState*  $\triangleq$  [  
     *proposals*       $\mapsto$  *proposal*,  
     *configurations*  $\mapsto$  *configuration*,  
     *targets*         $\mapsto$  *target*,  
     *masterships*    $\mapsto$  *mastership*,  
     *nodes*          $\mapsto$  *node*]

LOCAL *NextState*  $\triangleq$  [  
     *proposals*       $\mapsto$  *proposal'*,  
     *configurations*  $\mapsto$  *configuration'*,  
     *targets*         $\mapsto$  *target'*,  
     *masterships*    $\mapsto$  *mastership'*,  
     *nodes*          $\mapsto$  *node'*]

$\text{LOCAL } \text{Trace} \triangleq \text{INSTANCE } \text{Trace} \text{ WITH}$   
 $\text{Module} \leftarrow \text{"Proposal"},$   
 $\text{InitState} \leftarrow \text{InitState},$   
 $\text{NextState} \leftarrow \text{NextState},$   
 $\text{Enabled} \leftarrow \text{TraceProposal}$

---

Reconcile a proposal  
 $\text{ReconcileProposal}(n, i) \triangleq$   
 Only the master can process proposals for the target.  
 $\wedge \text{mastership.master} = n$   
 While in the Commit state, commit the proposed changes to the configuration.  
 $\wedge \vee \wedge \text{proposal}[i].\text{phase} = \text{ProposalCommit}$   
 $\wedge \vee \wedge \text{proposal}[i].\text{state} = \text{ProposalInProgress}$   
 Only commit the proposal if the prior proposal has already been committed.  
 $\wedge \text{configuration.committed.index} = i - 1$   
 For Change proposals validate the set of requested changes.  
 $\wedge \vee \wedge \text{proposal}[i].\text{type} = \text{ProposalChange}$   
 If all the change values are valid, record the changes required to roll  
 back the proposal and the index to which the rollback changes  
 will roll back the configuration.  
 $\wedge \vee \text{LET } \text{rollbackIndex} \triangleq \text{configuration.committed.revision}$   
 $\text{rollbackValues} \triangleq [p \in \text{DOMAIN } \text{proposal}[i].\text{change.values} \mapsto$   
 $\text{IF } p \in \text{DOMAIN } \text{configuration.committed.values} \text{ THEN}$   
 $\text{configuration.committed.values}[p]$   
 $\text{ELSE}$   
 $[index \mapsto 0, \text{value} \mapsto \text{Nil}]$   
 $\text{changeValues} \triangleq [p \in \text{DOMAIN } \text{proposal}[i].\text{change.values} \mapsto$   
 $\text{proposal}[i].\text{change.values}[p] \text{ @@ } [index \mapsto i]]$   
 $\text{IN } \wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } \begin{array}{l} \text{!.committed.index} = i, \\ \text{!.committed.revision} = i, \\ \text{!.committed.values} = \text{changeValues} \end{array}]$   
 $\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } \begin{array}{l} \text{![i].change} = [ \\ \text{index} \mapsto i, \\ \text{values} \mapsto \text{changeValues}], \\ \text{![i].rollback} = [ \\ \text{index} \mapsto \text{rollbackIndex}, \\ \text{values} \mapsto \text{rollbackValues}], \\ \text{![i].state} = \text{ProposalComplete} \end{array}]$   
 A proposal can fail validation at this point, in which case the proposal  
 is marked failed.  
 $\vee \wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } \text{![i].state} = \text{ProposalFailed}]$   
 $\wedge \text{UNCHANGED } \langle \text{configuration} \rangle$   
 For Rollback proposals, validate the rollback changes which are

proposal being rolled back.

$\vee \wedge \text{proposal}[i].\text{type} = \text{ProposalRollback}$

Rollbacks can only be performed on Change type proposals.

$\wedge \vee \wedge \text{proposal}[\text{proposal}[i].\text{rollback.index}].\text{type} = \text{ProposalChange}$

Only roll back the change if it's the latest change made  
to the configuration based on the configuration index.

$\wedge \vee \wedge \text{configuration.committed.revision} = \text{proposal}[i].\text{rollback.index}$

Record the changes required to roll back the target proposal and the index to  
which the configuration is being rolled back.

$\wedge \text{LET } \text{changeIndex} \triangleq \text{proposal}[\text{proposal}[i].\text{rollback.index}].\text{rollback.index}$   
 $\text{changeValues} \triangleq \text{proposal}[\text{proposal}[i].\text{rollback.index}].\text{rollback.values}$

Note: these two changes must be implemented as an atomic, idempotent update.

Implementations should check if the configuration has already been updated and  
skip the configuration update if the committed index is  $\geq$  the proposal index.

IN  $\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } !.\text{committed.index} = i,$   
 $!.\text{committed.revision} = \text{changeIndex},$   
 $!.\text{committed.values} = \text{changeValues}]$

$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![i].\text{change} = [$   
 $\text{index} \mapsto \text{changeIndex},$   
 $\text{values} \mapsto \text{changeValues}],$   
 $![i].\text{state} = \text{ProposalComplete}]$

If the Rollback target is not the most recent change to the configuration,  
fail validation for the proposal.

$\vee \wedge \text{configuration.committed.revision} \neq \text{proposal}[i].\text{rollback.index}$

Note: these two changes must be implemented as an atomic, idempotent update.

Implementations should check if the configuration has already been updated and  
skip the configuration update if the committed index is  $\geq$  the proposal index.

$\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } !.\text{committed.index} = i]$

$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![i].\text{state} = \text{ProposalFailed}]$

If a Rollback proposal is attempting to roll back another Rollback,  
fail validation for the proposal.

$\vee \wedge \text{proposal}[\text{proposal}[i].\text{rollback.index}].\text{type} = \text{ProposalRollback}$

Note: these two changes must be implemented as an atomic, idempotent update.

Implementations should check if the configuration has already been updated and  
skip the configuration update if the committed index is  $\geq$  the proposal index.

$\wedge \text{configuration}' = [\text{configuration} \text{ EXCEPT } !.\text{committed.index} = i]$

$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![i].\text{state} = \text{ProposalFailed}]$

$\wedge \text{UNCHANGED } \langle \text{target} \rangle$

Once the proposal is committed, update the configuration's commit index  
and move to the apply phase.

$\vee \wedge \text{proposal}[i].\text{state} = \text{ProposalComplete}$

$\wedge \text{proposal}' = [\text{proposal} \text{ EXCEPT } ![i].\text{phase} = \text{ProposalApply},$   
 $![i].\text{state} = \text{ProposalInProgress}]$

$\wedge \text{UNCHANGED } \langle \text{configuration}, \text{target} \rangle$

While in the Apply phase, apply the proposed changes to the target.

$\vee \wedge \text{proposal}[i].\text{phase} = \text{ProposalApply}$   
 $\wedge \text{configuration.applied.index} = i - 1$   
 $\wedge \text{proposal}[i].\text{state} = \text{ProposalInProgress}$   
 Process the proposal once the prior proposal has been applied.  
 $\wedge i - 1 \in \text{DOMAIN } \text{proposal} \Rightarrow$   
 $\vee \wedge \text{proposal}[i - 1].\text{phase} = \text{ProposalCommit}$   
 $\wedge \text{proposal}[i - 1].\text{state} = \text{ProposalFailed}$   
 $\vee \wedge \text{proposal}[i - 1].\text{phase} = \text{ProposalApply}$   
 $\wedge \text{proposal}[i - 1].\text{state} \in \{\text{ProposalComplete}, \text{ProposalFailed}\}$   
 Verify the applied term is the current *mastership* term to ensure the  
 configuration has been synchronized following restarts.  
 $\wedge \text{configuration.applied.term} = \text{mastership.term}$   
 Verify the node's connection to the target.  
 $\wedge \text{node}[n].\text{connected}$   
 $\wedge \text{target.running}$   
 Model successful and failed target update requests.  
 $\wedge \vee \wedge \text{target}' = [\text{target } \text{EXCEPT } !.\text{values} = \text{proposal}[i].\text{change.values}]$   
 Note: these two changes must be implemented as an atomic, idempotent update.  
 Implementations should check if the configuration has already been updated and  
 skip the configuration update if the applied index is  $\geq$  the proposal index.  
 $\wedge \text{LET } \text{index} \triangleq \text{proposal}[i].\text{change.index}$   
 $\text{values} \triangleq \text{proposal}[i].\text{change.values} @@ \text{configuration.applied.values}$   
 $\text{IN } \text{configuration}' = [\text{configuration } \text{EXCEPT } !.\text{applied.index} = i,$   
 $\text{!.applied.revision} = \text{index},$   
 $\text{!.applied.values} = \text{values}]$   
 $\wedge \text{proposal}' = [\text{proposal } \text{EXCEPT } ![i].\text{state} = \text{ProposalComplete}]$   
 If the proposal could not be applied, update the configuration's applied index  
 and mark the proposal Failed.  
 Note: these two changes must be implemented as an atomic, idempotent update.  
 Implementations should check if the configuration has already been updated and  
 skip the configuration update if the applied index is  $\geq$  the proposal index.  
 $\vee \wedge \text{configuration}' = [\text{configuration } \text{EXCEPT } !.\text{applied.index} = i]$   
 $\wedge \text{proposal}' = [\text{proposal } \text{EXCEPT } ![i].\text{state} = \text{ProposalFailed}]$   
 $\wedge \text{UNCHANGED } \langle \text{target} \rangle$   
 $\wedge \text{UNCHANGED } \langle \text{mastership}, \text{node} \rangle$

---

Formal specification, constraints, and theorems.

$\text{InitProposal} \triangleq$   
 $\wedge \text{proposal} = [$   
 $i \in \{\} \mapsto [$   
 $\text{type} \mapsto \text{ProposalChange},$   
 $\text{change} \mapsto [$   
 $\text{index} \mapsto 0,$

$$\begin{aligned}
& \text{values} \mapsto [p \in \{\} \mapsto [\text{index} \mapsto 0, \text{value} \mapsto \text{Nil}, \text{delete} \mapsto \text{FALSE}]]], \\
& \text{rollback} \mapsto [ \\
& \quad \text{index} \mapsto 0, \\
& \quad \text{values} \mapsto [p \in \{\} \mapsto [\text{index} \mapsto 0, \text{value} \mapsto \text{Nil}, \text{delete} \mapsto \text{FALSE}]]], \\
& \text{phase} \mapsto \text{ProposalCommit}, \\
& \text{state} \mapsto \text{ProposalInProgress}] \\
& \wedge \text{Trace!Init} \\
& \text{NextProposal} \triangleq \\
& \quad \forall \exists n \in \text{Nodes} : \\
& \quad \quad \exists i \in \text{DOMAIN } \text{proposal} : \\
& \quad \quad \text{Trace!Step}(\text{ReconcileProposal}(n, i), [\text{node} \mapsto n, \text{index} \mapsto i])
\end{aligned}$$


---

\ \* Modification History  
\ \* Last modified *Fri Apr 21 19:15:11 PDT 2023* by *jhalterm*  
\ \* Last modified *Mon Feb 21 01:24:12 PST 2022* by *jordanhalterm*  
\ \* Created *Sun Feb 20 10:07:16 PST 2022* by *jordanhalterm*