
MODULE *Transaction*

INSTANCE *Naturals*

INSTANCE *FiniteSets*

INSTANCE *Sequences*

INSTANCE *TLC*

An empty constant

CONSTANT *Nil*

Transaction type constants

CONSTANTS

Change,
Rollback

Transaction isolation constants

CONSTANTS

ReadCommitted,
Serializable

Phase constants

CONSTANTS

Initialize,
Validate,
Abort,
Commit,
Apply

Phase \triangleq

$\{$ *Initialize*,
Validate,
Abort,
Commit,
Apply $\}$

Status constants

CONSTANTS

InProgress,
Complete,
Failed

State \triangleq

$\{$ *InProgress*,
Complete,
 $\}$

Failed}

State constants

CONSTANTS

Pending,
Validated,
Committed,
Applied,
Aborted

Status \triangleq

{*Pending*,
Validated,
Committed,
Applied,
Aborted}

CONSTANTS

Valid,
Invalid

CONSTANTS

Success,
Failure

The set of all nodes

CONSTANT *Node*

Target is the set of all targets and their possible paths and values.

Example:

Target \triangleq

[*target1* \mapsto
[*persistent* \mapsto FALSE, *values* \mapsto [
 path1 \mapsto {"*value1*", "*value2*"},
 path2 \mapsto {"*value2*", "*value3*" }]],
target2 \mapsto
[*persistent* \mapsto TRUE, *values* \mapsto [
 path2 \mapsto {"*value3*", "*value4*"},
 path3 \mapsto {"*value4*", "*value5*" }]]]

CONSTANT *Target*

Empty \triangleq [*p* \in {}] \mapsto [*value* \mapsto *Nil*, *delete* \mapsto FALSE]]

A transaction log. Transactions may either request a set of changes to a set of targets or rollback a prior change.

VARIABLE *transaction*

A record of per-target proposals
 VARIABLE *proposal*

A record of per-target configurations
 VARIABLE *configuration*

A record of target states
 VARIABLE *target*

A record of target masterships
 VARIABLE *mastership*

$Test \triangleq$ INSTANCE *Test* WITH
 File \leftarrow "Transaction.log",
 CurrState \leftarrow [
 transactions \mapsto *transaction*,
 proposals \mapsto *proposal*,
 configuration \mapsto *configuration*,
 mastership \mapsto *mastership*,
 target \mapsto *target*],
 SuccState \leftarrow [
 transactions \mapsto *transaction'*,
 proposals \mapsto *proposal'*,
 configuration \mapsto *configuration'*,
 mastership \mapsto *mastership'*,
 target \mapsto *target'*]

This section models configuration changes and rollbacks. Changes are appended to the transaction log and processed asynchronously.

Add a set of changes 'c' to the transaction log

$RequestChange(i, c) \triangleq$
 $\wedge i = Len(transaction) + 1$
 $\wedge \exists isolation \in \{ReadCommitted, Serializable\} :$
 $\wedge transaction' = transaction @@ (i :> [type \mapsto Change,$
 $isolation \mapsto isolation,$
 $change \mapsto c,$
 $targets \mapsto \{\},$
 $phase \mapsto Initialize,$
 $state \mapsto InProgress,$
 $status \mapsto Pending])$
 $\wedge UNCHANGED \langle proposal, configuration, mastership, target \rangle$

Add a rollback of transaction 't' to the transaction log

$RequestRollback(i, j) \triangleq$
 $\wedge i = Len(transaction) + 1$

$$\begin{aligned}
& \wedge \exists \text{isolation} \in \{\text{ReadCommitted}, \text{Serializable}\} : \\
& \quad \wedge \text{transaction}' = \text{transaction} @@ (i :> [type \mapsto \text{Rollback}, \\
& \quad \quad \quad \text{isolation} \mapsto \text{isolation}, \\
& \quad \quad \quad \text{rollback} \mapsto j, \\
& \quad \quad \quad \text{targets} \mapsto \{\}, \\
& \quad \quad \quad \text{phase} \mapsto \text{Initialize}, \\
& \quad \quad \quad \text{state} \mapsto \text{InProgress}, \\
& \quad \quad \quad \text{status} \mapsto \text{Pending}]) \\
& \wedge \text{UNCHANGED} \langle \text{proposal}, \text{configuration}, \text{mastership}, \text{target} \rangle
\end{aligned}$$

This section models the *Transaction* log reconciler.

Transactions come in two flavors: - *Change* transactions contain a set of changes to be applied to a set of targets - *Rollback* transactions reference a prior change transaction to be reverted to the previous state

Transactions proceed through a series of phases:

- * *Initialize* - create and link Proposals
- * *Validate* - validate changes and rollbacks
- * *Commit* - commit changes to Configurations
- * *Apply* - commit changes to Targets

Reconcile a transaction

$$\text{ReconcileTransaction}(n, i) \triangleq$$

Initialize is the only transaction phase that's globally serialized.

While in the Initializing phase, the reconciler checks whether the prior transaction has been Initialized before creating Proposals in the *Initialize* phase. Once all of the transaction's proposals have been Initialized, the transaction will be marked Initialized. If any proposal is *Failed*, the transaction will be marked *Failed* as well.

$$\wedge \vee \wedge \text{transaction}[i].\text{phase} = \text{Initialize}$$

$$\wedge \vee \wedge \text{transaction}[i].\text{state} = \text{InProgress}$$

All prior transaction must be initialized before proceeding to initialize this transaction.

$$\wedge \neg \exists j \in \text{DOMAIN } \text{transaction} :$$

$$\wedge j < i$$

$$\wedge \text{transaction}[j].\text{phase} = \text{Initialize}$$

$$\wedge \text{transaction}[j].\text{state} = \text{InProgress}$$

If the transaction's targets are not yet set, create proposals and add targets to the transaction state.

$$\wedge \vee \wedge \text{transaction}[i].\text{targets} = \{\}$$

If the transaction is a change, the targets are taken from the change values.

$$\wedge \vee \wedge \text{transaction}[i].\text{type} = \text{Change}$$

$$\wedge \text{transaction}' = [\text{transaction } \text{EXCEPT } ![i].\text{targets} = \text{DOMAIN } \text{transaction}[i].\text{change}]$$

$$\wedge \text{proposal}' = [t \in \text{DOMAIN } \text{proposal} \mapsto$$

IF $t \in \text{DOMAIN } \text{transaction}[i].\text{change}$ THEN
 $\text{proposal}[t] @@ (i := [type \mapsto \text{Change},$
 $\text{change} \mapsto$
 $[index \mapsto i,$
 $\text{values} \mapsto \text{transaction}[i].\text{change}[t]],$
 $\text{rollback} \mapsto$
 $[index \mapsto 0,$
 $\text{values} \mapsto \text{Empty}],$
 $\text{dependency} \mapsto [index \mapsto 0],$
 $\text{phase} \mapsto \text{Initialize},$
 $\text{state} \mapsto \text{InProgress})$

 ELSE
 $\text{proposal}[t]$
 If the transaction is a rollback, the targets affected are
 the targets of the change transaction being rolled back.
 $\vee \wedge \text{transaction}[i].\text{type} = \text{Rollback}$
 If the rollback index is a valid *Change* transaction,
 initialize proposals for all of the *Change* targets.
 $\wedge \vee \wedge \text{transaction}[i].\text{rollback} \in \text{DOMAIN } \text{transaction}$
 $\wedge \text{transaction}[\text{transaction}[i].\text{rollback}].\text{type} = \text{Change}$
 $\wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{targets} =$
 $\text{DOMAIN } \text{transaction}[\text{transaction}[i].\text{rollback}].\text{change}]$
 $\wedge \text{proposal}' = [t \in \text{DOMAIN } \text{proposal} \mapsto$
 IF $t \in \text{DOMAIN } \text{transaction}[\text{transaction}[i].\text{rollback}].\text{change}$ THEN
 $\text{proposal}[t] @@ (i := [type \mapsto \text{Rollback},$
 $\text{change} \mapsto$
 $[index \mapsto 0,$
 $\text{values} \mapsto \text{Empty}],$
 $\text{rollback} \mapsto$
 $[index \mapsto \text{transaction}[i].\text{rollback},$
 $\text{values} \mapsto \text{Empty}],$
 $\text{dependency} \mapsto [index \mapsto 0],$
 $\text{phase} \mapsto \text{Initialize},$
 $\text{state} \mapsto \text{InProgress})$

 ELSE
 $\text{proposal}[t]$
 If the rollback index is not a valid *Change* transaction
 fail the *Rollback* transaction.
 $\vee \wedge \vee \wedge \text{transaction}[i].\text{rollback} \in \text{DOMAIN } \text{transaction}$
 $\wedge \text{transaction}[\text{transaction}[i].\text{rollback}].\text{type} = \text{Rollback}$
 $\vee \text{transaction}[i].\text{rollback} \notin \text{DOMAIN } \text{transaction}$
 $\wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{state} = \text{Failed}]$
 $\wedge \text{UNCHANGED } \langle \text{proposal} \rangle$
 If the transaction's proposals have been initialized, check proposals
 for completion or failures.

If all proposals have been *Complete*, mark the transaction *Complete*.

$$\wedge proposal[t][i].phase = Initialize$$
$$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Complete]$$

If any proposal has been *Failed*, mark the transaction *Failed*.

$$\wedge proposal[t][i].phase = Initialize$$
$$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Failed]$$

If any of the transaction's proposals depend on a *Serializable* transaction,

$$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Complete,$$

$![i].status = Validated]$

\wedge UNCHANGED $\langle proposal \rangle$

If any proposal has been *Failed*, mark the transaction *Failed*.

$\vee \wedge \exists t \in transaction[i].targets :$

$\wedge proposal[t][i].phase = Validate$

$\wedge proposal[t][i].state = Failed$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Failed]$

\wedge UNCHANGED $\langle proposal \rangle$

Once the transaction has been *Validated*, proceed to the *Commit* phase.

If any of the transaction's proposals depend on a *Serializable* transaction, verify the dependency has been *Committed* to preserve serializability before moving the transaction to the *Commit* phase.

$\vee \wedge transaction[i].state = Complete$

$\wedge \forall t \in transaction[i].targets :$

$\wedge proposal[t][i].dependency.index \in \text{DOMAIN } transaction$

$\wedge transaction[proposal[t][i].dependency.index].isolation = Serializable$

$\Rightarrow transaction[proposal[t][i].dependency.index].status \in \{Committed, Applied, Aborted\}$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Commit,$

$![i].state = InProgress]$

\wedge UNCHANGED $\langle proposal \rangle$

If the transaction failed validation, proceed to the *Abort* phase

to ensure indexes are still updated for the target configurations.

$\vee \wedge transaction[i].state = Failed$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Abort,$

$![i].state = InProgress]$

\wedge UNCHANGED $\langle proposal \rangle$

$\vee \wedge transaction[i].phase = Commit$

$\wedge \vee \wedge transaction[i].state = InProgress$

Move the transaction's proposals to the Committing state

$\wedge \vee \wedge \exists t \in transaction[i].targets :$

$\wedge proposal[t][i].phase \neq Commit$

$\wedge proposal' = [proposal \text{ EXCEPT } ![t] =$

$[proposal[t] \text{ EXCEPT } ![i].phase = Commit,$

$![i].state = InProgress]]$

\wedge UNCHANGED $\langle transaction \rangle$

If all proposals have been *Complete*, mark the transaction *Complete*.

$\vee \wedge \forall t \in transaction[i].targets :$

$\wedge proposal[t][i].phase = Commit$

$\wedge proposal[t][i].state = Complete$

$\wedge transaction' = [transaction \text{ EXCEPT } ![i].state = Complete,$

$![i].status = Committed]$

\wedge UNCHANGED $\langle proposal \rangle$

Once the transaction has been *Committed*, proceed to the *Apply* phase.

If any of the transaction's proposals depend on a *Serializable* transaction, verify the dependency has been *Applied* to preserve serializability before

$$\vee \wedge transaction[i].state = Complete$$
$$\wedge proposal[t][i].dependenc$$
$$\wedge transaction[proposal[t][i].dependency.index].isolation = Serializable$$
$$\wedge transaction' = [transaction \text{ EXCEPT } ![i].phase = Apply,$$
 \wedge UNCHANGED $\langle proposal \rangle$
$$\wedge transaction[i].state = InProgress$$
$$\wedge \vee \wedge \exists t \in transaction[i].targets :$$
$$\wedge proposal' = [proposal \text{ EXCEPT } ![t] =$$
$$![i].state = InProgress]]$$
$$\vee \wedge \forall t \in transaction[i].targets :$$
$$\wedge proposal[t][i].state = Complete$$
$$![i].status = Applied]$$
$$\forall \wedge \exists t \in transaction[i].targets :$$
$$\wedge proposal[t][i].state = Failed$$
 \wedge UNCHANGED $\langle proposal \rangle$
$$\vee \wedge transaction[i].phase = Abort$$

Move the transaction's proposals to the Aborting state

$$\wedge \text{proposal}[t][i].\text{phase} \neq \text{Abort}$$
$$[proposal[t] \text{ EXCEPT } ! [i].phase = Abort,$$
 \wedge UNCHANGED $\langle transaction \rangle$
$$\vee \wedge \forall t \in transaction[i].targets :$$

8

$$\begin{aligned}
& \wedge \text{proposal}[t][i].\text{state} = \text{Complete} \\
& \wedge \text{transaction}' = [\text{transaction} \text{ EXCEPT } ![i].\text{state} = \text{Complete}, \\
& \hspace{15em} ![i].\text{status} = \text{Aborted}] \\
& \wedge \text{UNCHANGED } \langle \text{proposal} \rangle \\
& \wedge \text{UNCHANGED } \langle \text{configuration}, \text{mastership}, \text{target} \rangle
\end{aligned}$$
